

Hypothesis Testing

Wolfgang Huber, EMBL



Das Orakel zu Delphi.

fineart
america

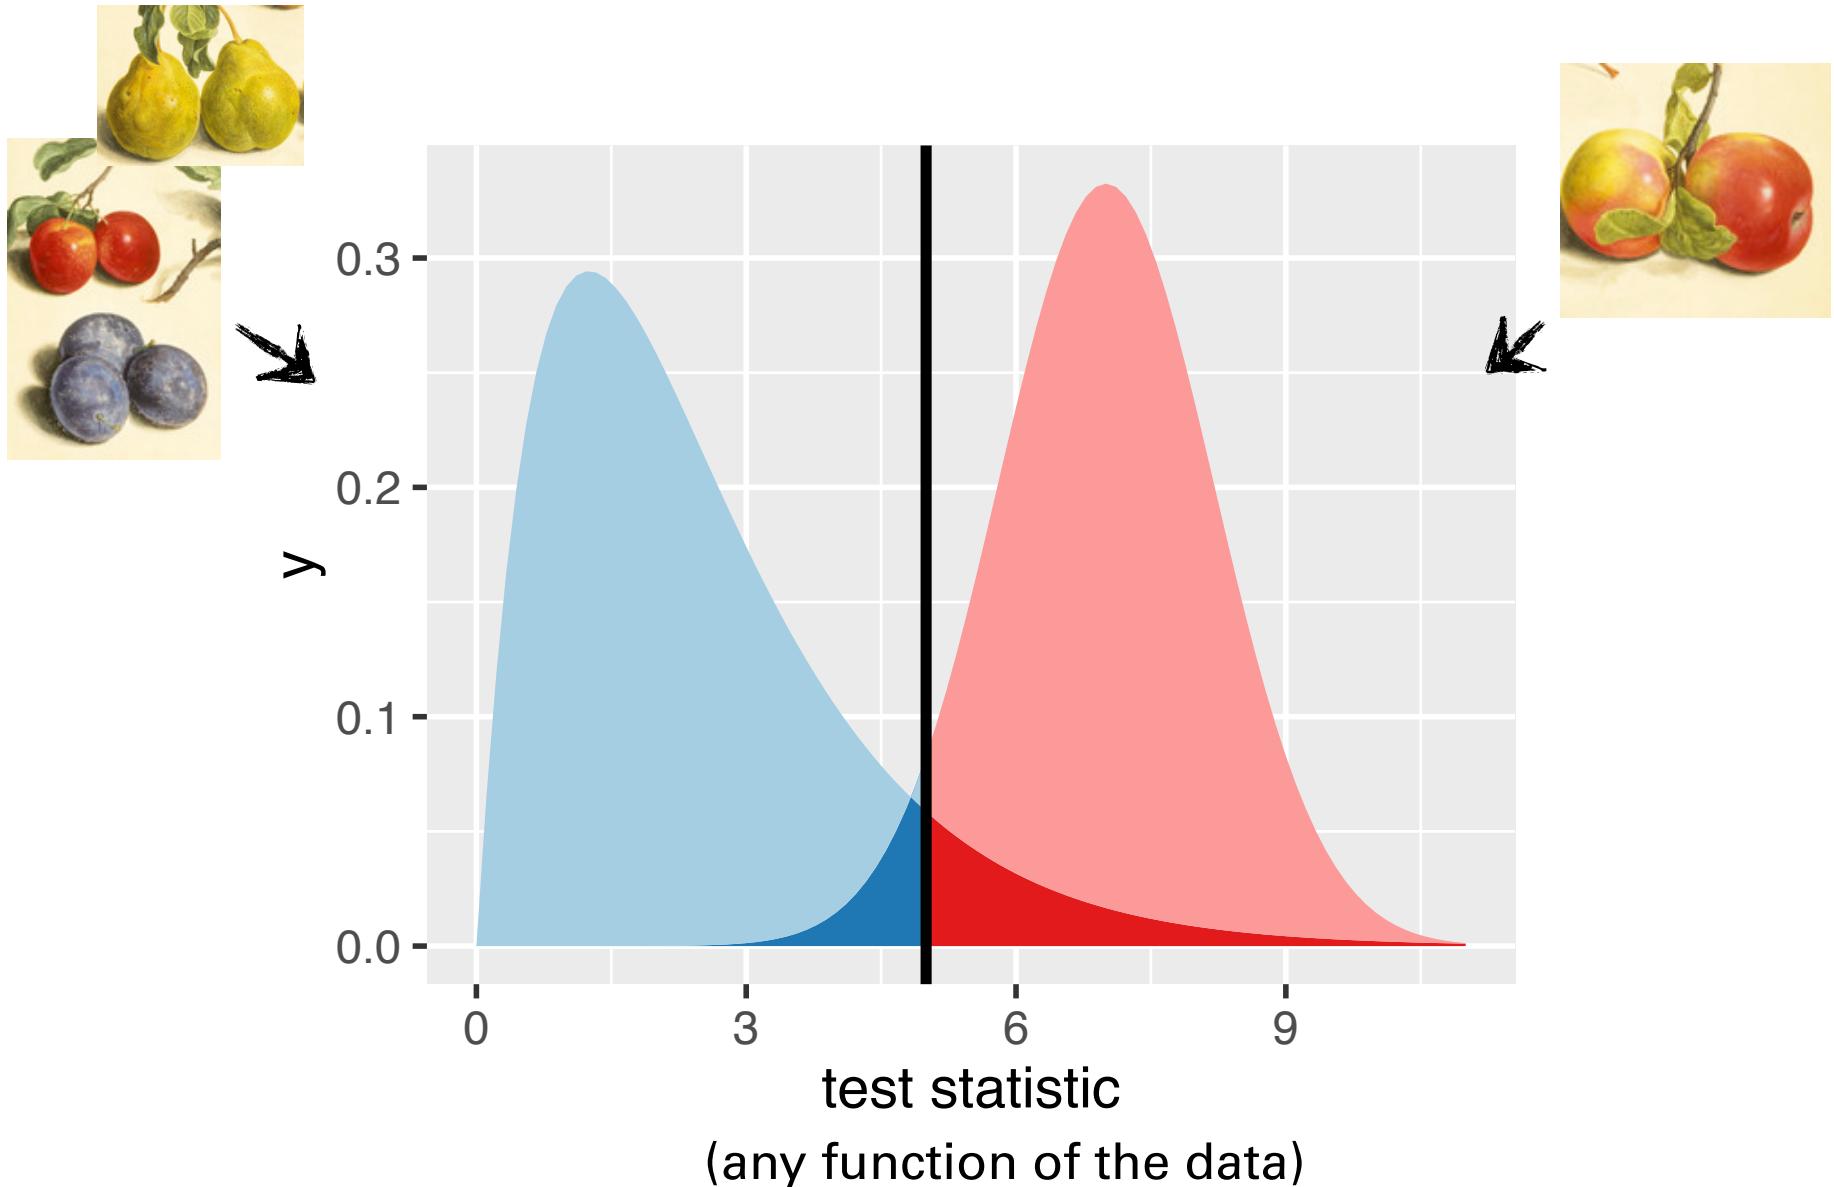
Aims for this Lecture

Understand the basic principles of hypothesis testing, its pitfalls, strengths, use cases and limitations

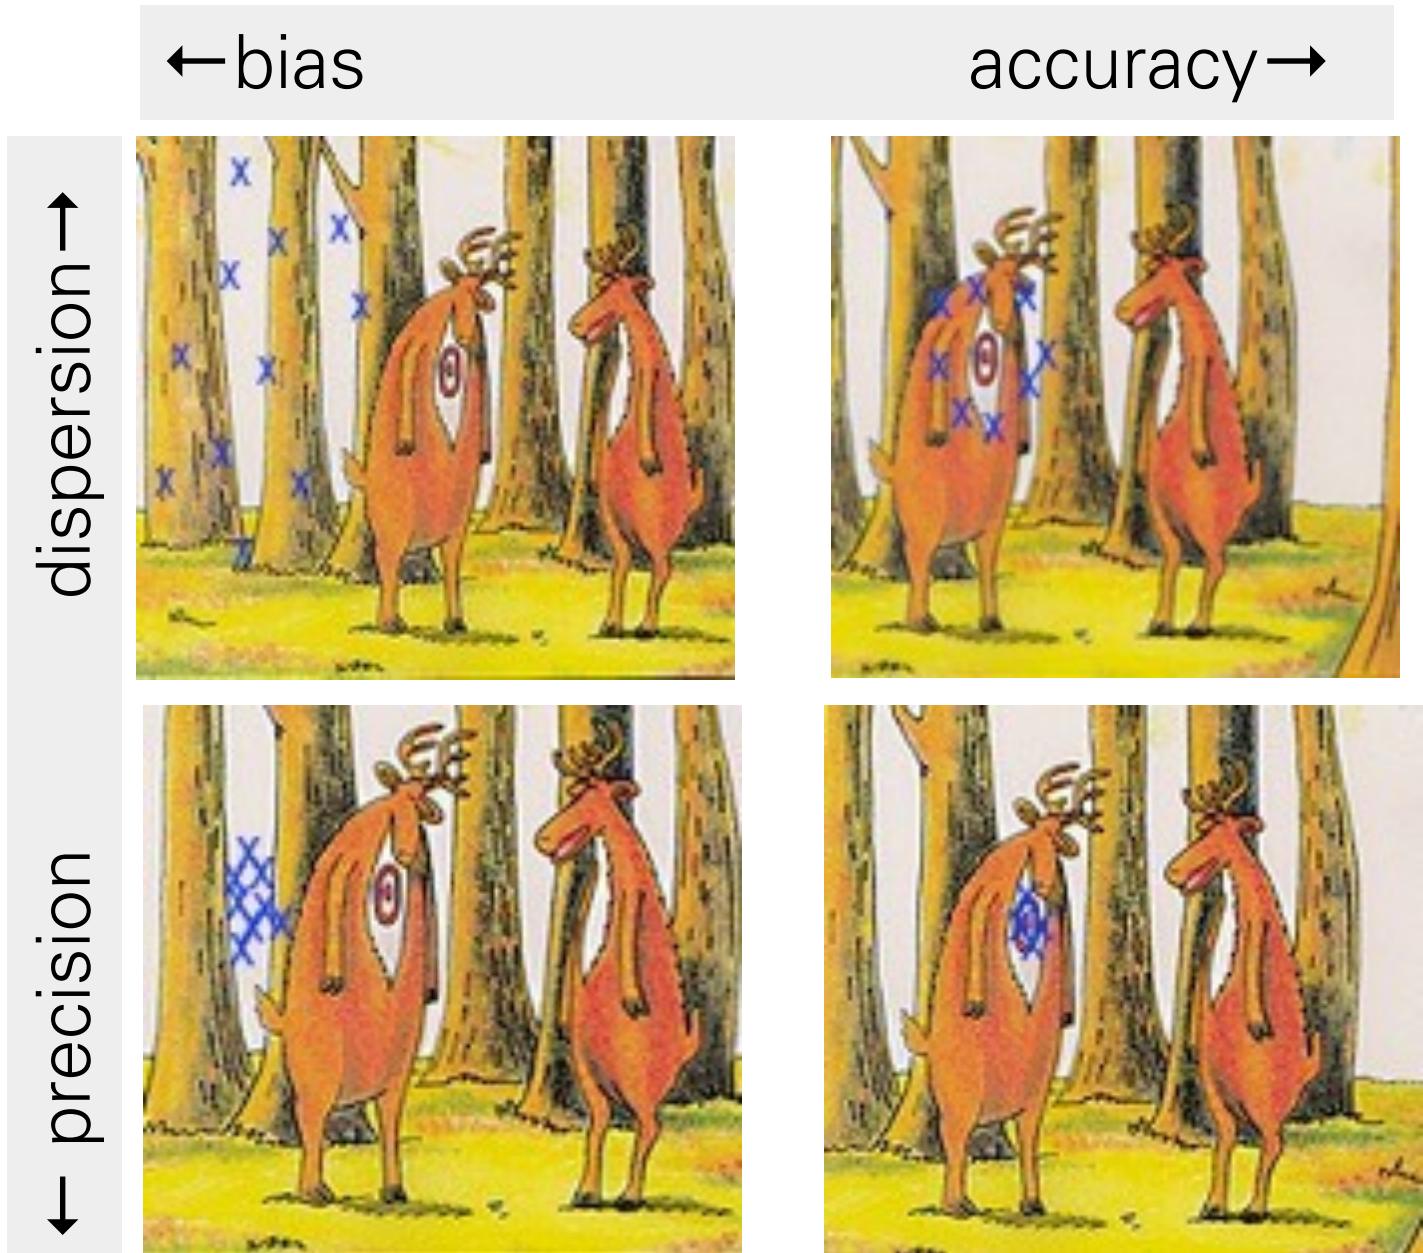
What changes when we go from single to multiple testing?

False discovery rates, p-value ‘adjustments’, filtering and weighting

Testing vs Classification



Accuracy vs Precision - Bias vs Variance



How to make rational decisions based on noisy, finite data

Prototypical examples:

- Testing efficacy of a drug on people
 - lack of complete experimental control
 - finite sample size
 - Effect of fertilizer, genetic variants, ... on phenotype of plants in an outdoors field trial
 - Lady testing tea, clairvoyant, telepath, ...
 - Toxicology
- +**: No understanding of mechanism involved/needed/desired
- : Wouldn't we want to use any available understanding or 'priors'?

Example

Toss a coin a number of times ⇒

If the coin is fair, then heads should appear half of the time (roughly).



But what is “roughly”? We use combinatorics / probability theory to quantify this.

Suppose we flipped the coin 100 times and got 59 heads. Is this ‘significant’?

Binomial Distribution

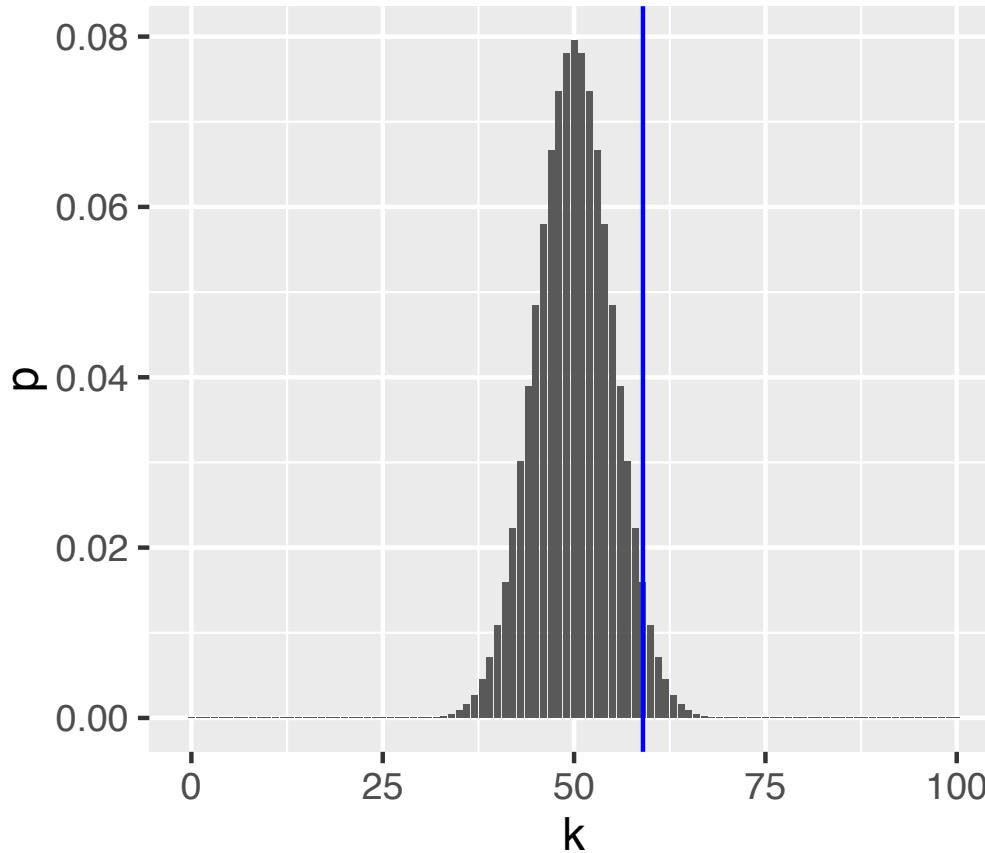


Figure 6.3: The binomial distribution for the parameters $n = 100$ and $p = 0.5$,

$$P(K = k | n, p) = \binom{n}{k} p^k (1 - p)^{n-k}$$

Rejection Region

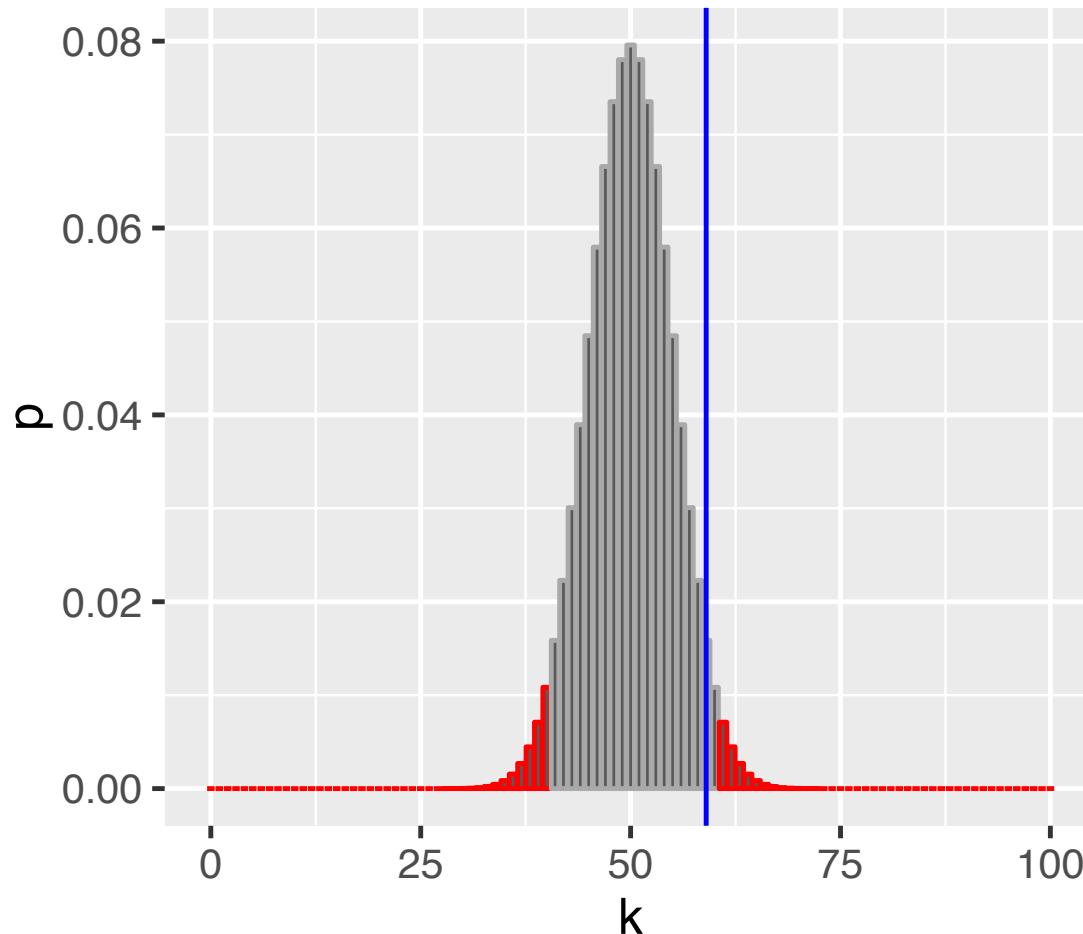


Figure 6.5: As Figure 6.3, with rejection region (red) whose total area is $\alpha = 0.05$.

Questions

- Does the fact that we don't reject the null hypothesis mean that the coin is fair?
- Would we have a better chance of detecting an unfair coin if we did more coin tosses? How many?
- If we repeated the whole procedure and again tossed the coin 100 times, might we **then** reject the null hypothesis?
- Our rejection region is asymmetric - its left part ends with 40, while its right part starts with 61. Why is that? Which other ways of defining the rejection region might be useful?

The Five Steps of Hypothesis Testing

Choose an experimental design and a data summary function for the effect that you are interested in: the **test statistic**

Set up a **null hypothesis**: a simple, computationally tractable model of reality that lets you compute the null distribution of the test statistic, i.e. all its possible outcomes and each of their probabilities.

Decide on the **rejection region**, i.e., a subset of possible outcomes whose total probability is small (**significance level**).

Do the experiment, collect data, compute the test statistic.

Make a **decision**: reject null hypothesis if the test statistic is in the rejection region.



The Five Steps of Hypothesis Testing

Choose an experimental design and a data summary function for the effect that you are interested in: the **test statistic**

Set up a

~~null hypothesis: a simple, computationally tractable~~

model of

the

test stati

This is the idealised scenario,
“orthodoxy”

probabil

Reality, esp. in retrospective ‘data-
mining’ can be quite different.

Decide

outcom

(signific

Do the e

compute the test statistic.

Make a **decision**: reject null hypothesis
if the test statistic is in the rejection region.



Examples of null hypotheses:

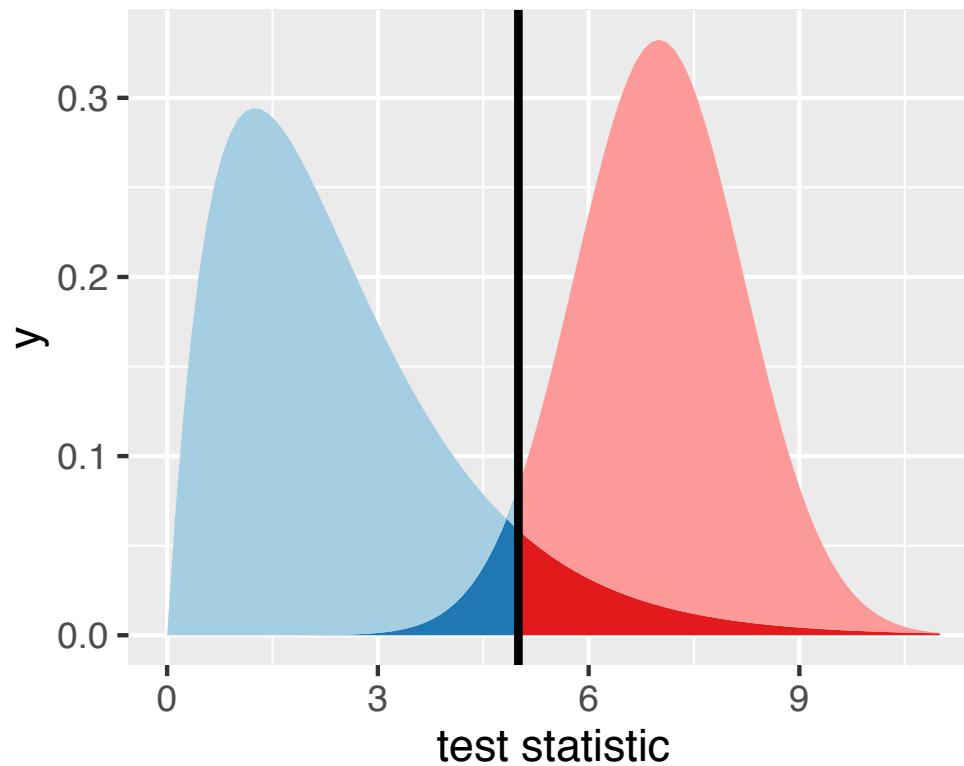
- The coin is fair
- The new drug is no better or worse than a placebo
- The effect of that CRISPR-edit on my cells is no different than that of a negative control treatment

These are not null hypotheses:

- The number of heads and tails were the same
- The coin is not fair
- The drug should not be approved

Types of Error in Testing

Test vs reality	Null hypothesis is true	... is false
Reject null hypothesis	Type I error (false positive)	True positive
Do not reject	True negative	Type II error (false negative)



Parametric Theory vs Simulation

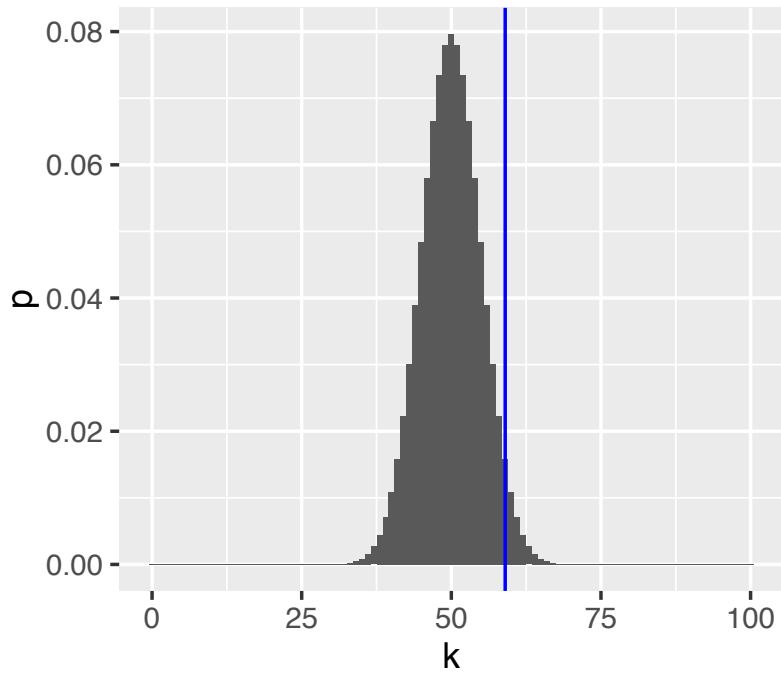


Figure 6.3: The binomial distribution for the parameters $n = 100$ and $p = 0.5$, according to Equation (6.1).

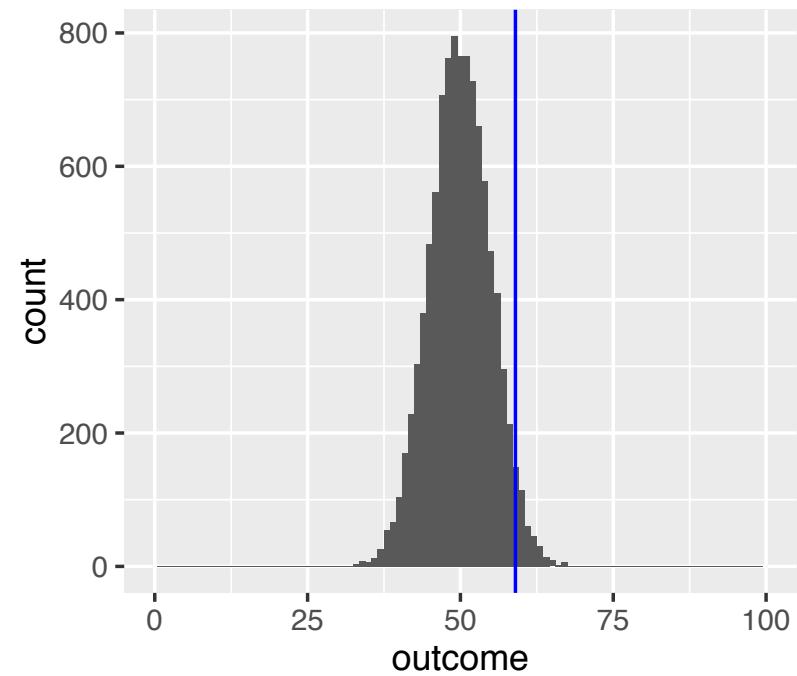


Figure 6.4: An approximation of the binomial distribution from 10^4 simulations (same parameters as Figure 6.3).

Parametric Theory vs Simulation

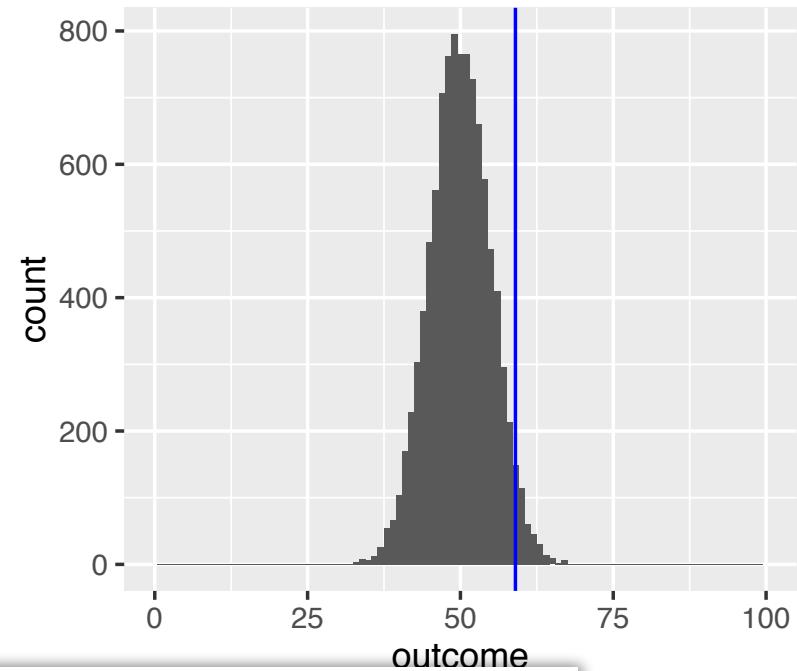
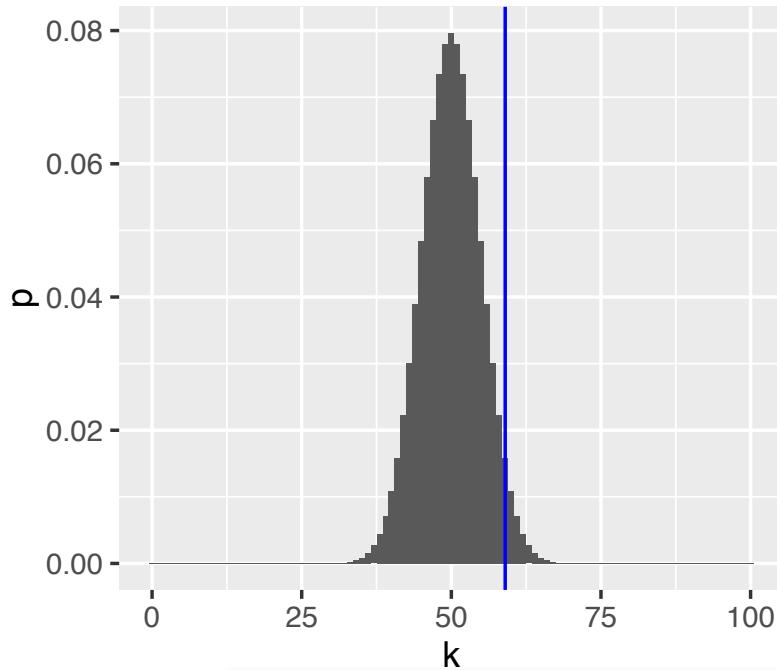


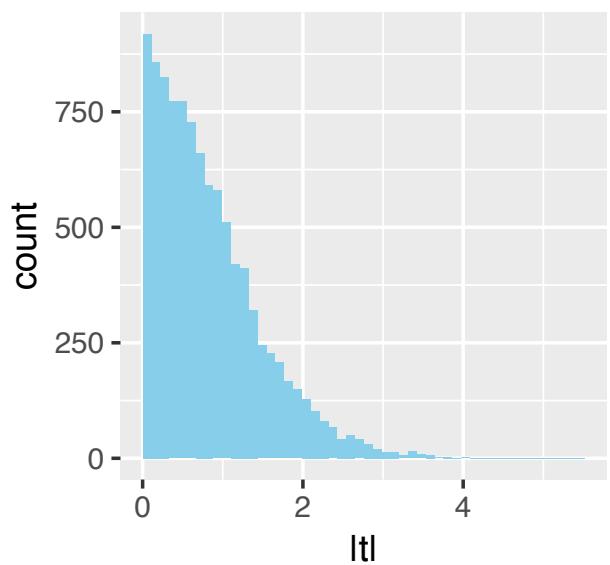
Figure 6.3: The parameter according to

Q:

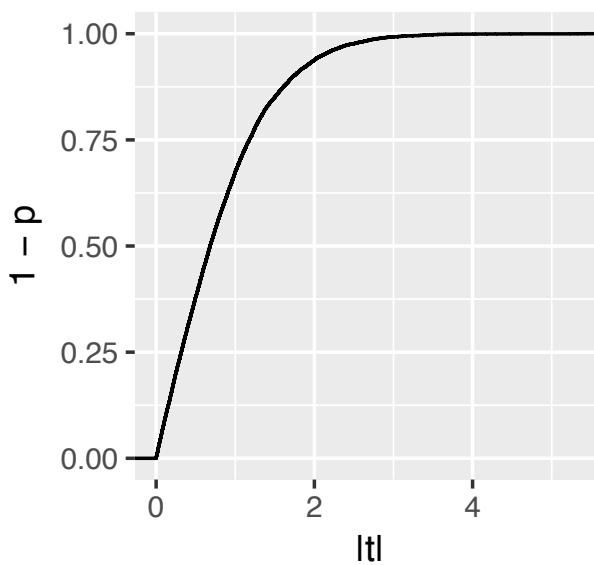
Discuss pros and contras for each

of the simulations).

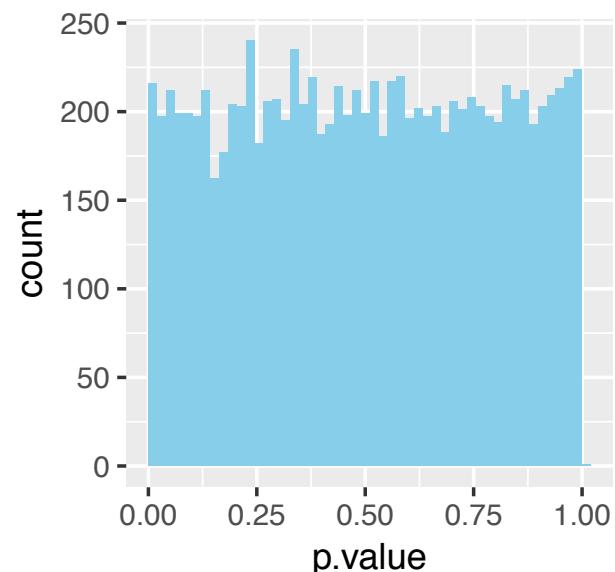
p-Values as Random Variables



test statistic

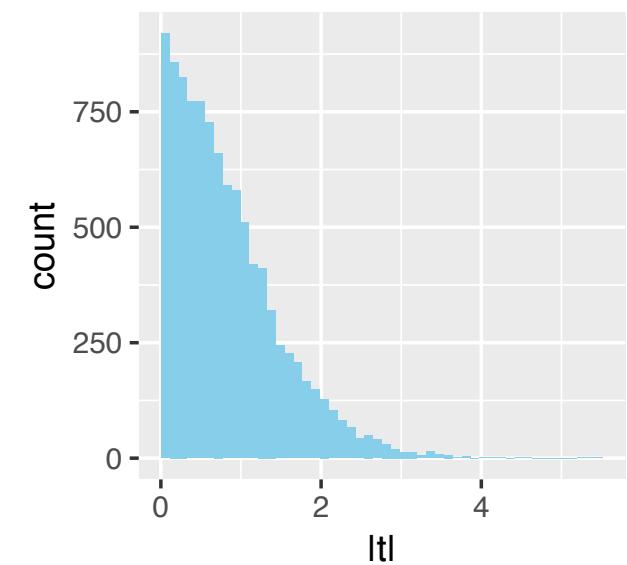


distribution
function

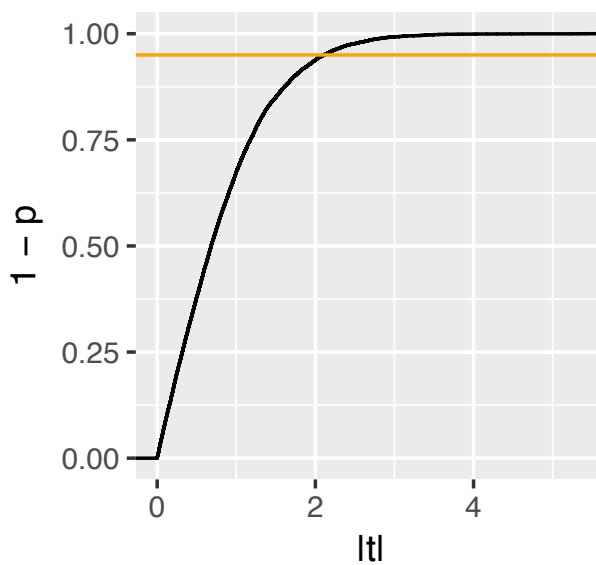


p-value

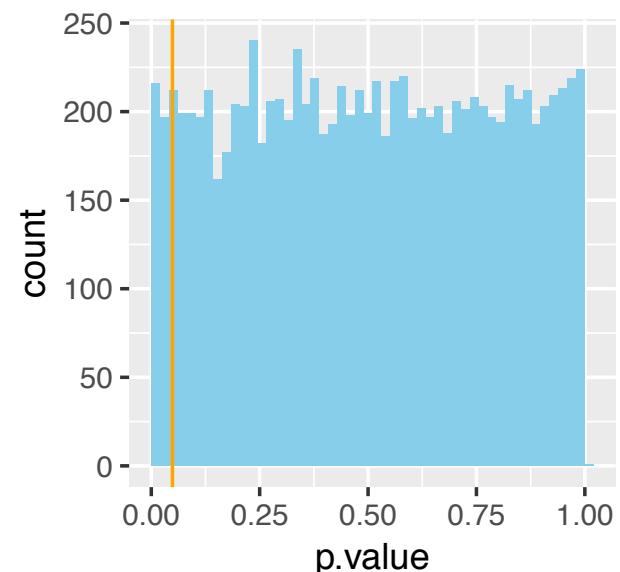
p-Values as Random Variables



test statistic

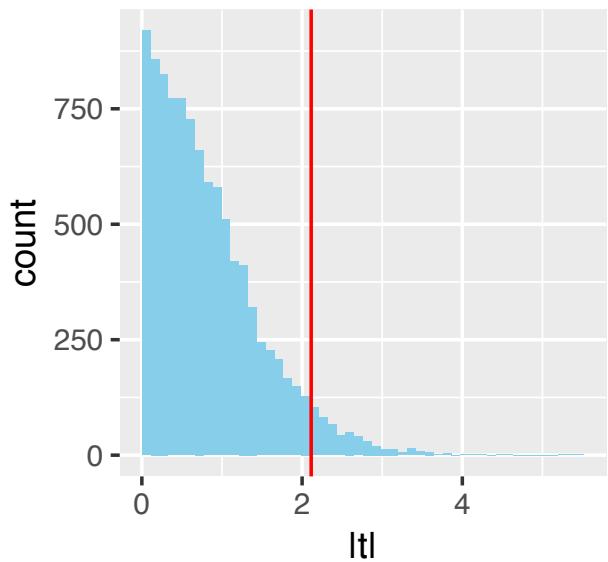
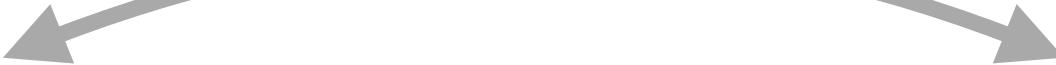


distribution
function

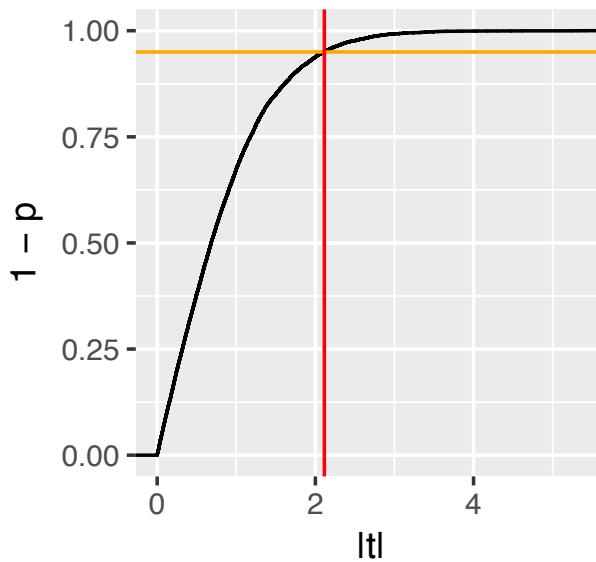


p-value

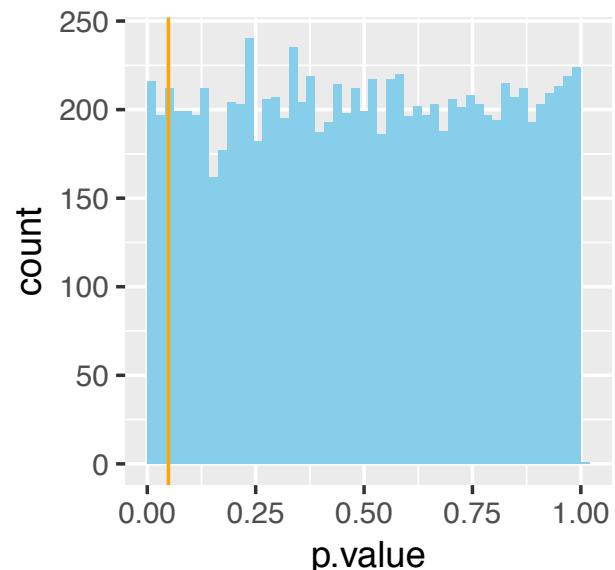
p-Values as Random Variables



test statistic



distribution
function



p-value

The Test Statistic

Suppose we observed 50 tails in a row, and then 50 heads in a row. Is this a perfectly fair coin?

We could use a different test statistic: number of times we see two tails in a row

Is this statistic generally and always preferable?

Power

There can be several test statistics, with different power, for different types of alternative

continuous data: the t-statistic

$$t = c \frac{m_1 - m_2}{s}$$

- Can also be adapted to one group only
- Relation to z-score

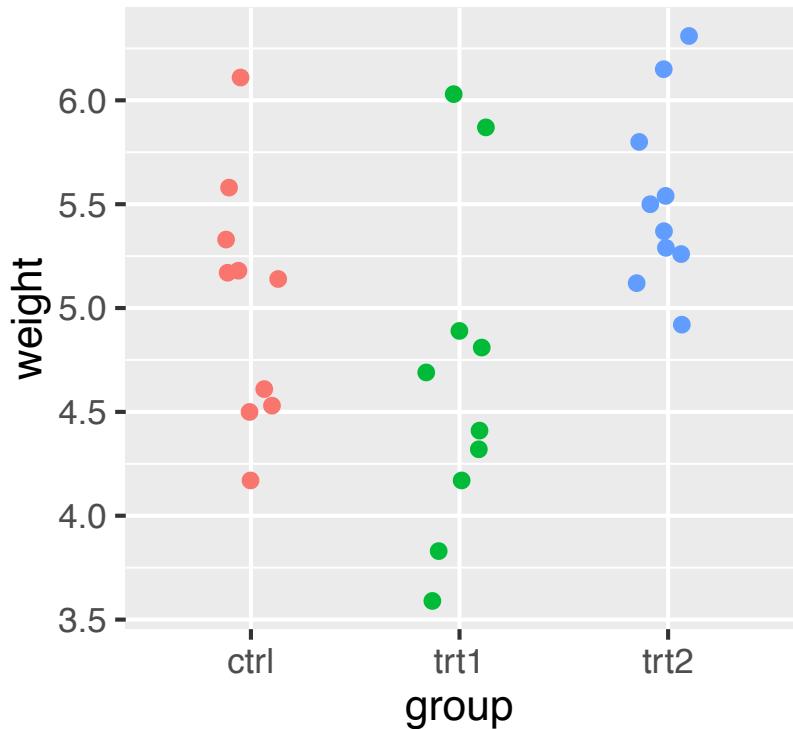


Figure 6.7: The PlantGrowth data.

$$m_g = \frac{1}{n_g} \sum_{i=1}^{n_g} x_{g,i} \quad g = 1, 2$$

$$s^2 = \frac{1}{n_1 + n_2 - 2} \left(\sum_{i=1}^{n_1} (x_{1,i} - m_1)^2 + \sum_{j=1}^{n_2} (x_{2,j} - m_2)^2 \right)$$

$$c = \sqrt{\frac{n_1 n_2}{n_1 + n_2}}.$$

t- (and |t|-) Distribution

If the data are identically normal distributed and independent, then under H_0 , t follows a ' t -distribution' with parameter n_1+n_2 (a.k.a. degrees of freedom)

t- (and $|t|$ -) Distribution

If the data are identically normally distributed and independent, then under H_0 , t follows a 't-distribution' with parameter n_1+n_2 (a.k.a. degrees of freedom)

Q:

How does the distribution of $|t|$ look?

t- (and |t|-) Distribution

If the data are identically normal distributed and independent, then under H_0 , t follows a 't-distribution' with parameter n_1+n_2 (a.k.a. degrees of freedom)

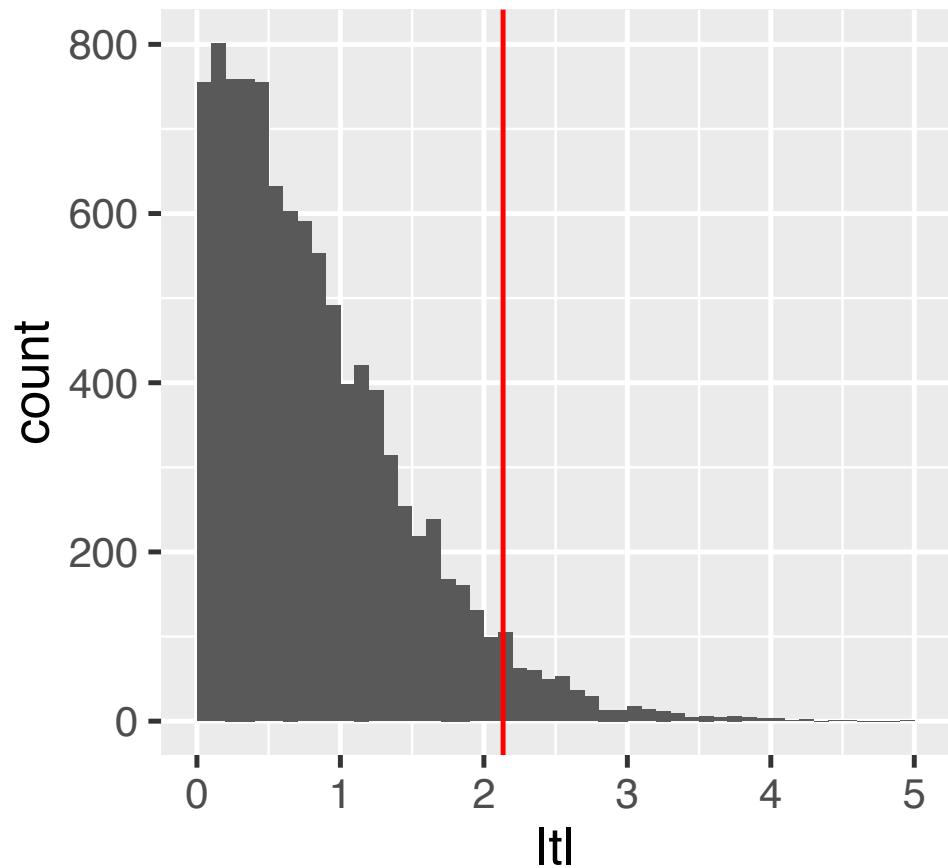


Figure 6.8: The null distribution of the (absolute) t -statistic determined by simulations – namely, by random permutations of the group labels.

Comments and Pitfalls

The proof that the t -statistic follows a t -distribution assumes that observations are independent and follow a normal distribution: a sufficient, but not necessary, condition

Comments and Pitfalls

The proof that the t -statistic follows a t -distribution assumes that observations are independent and follow a normal distribution: a sufficient, but not necessary, condition

[Deviation from normality](#) (heavier tails): test typically maintains type-I error control, but no longer has optimal power.

Comments and Pitfalls

The proof that the t -statistic follows a t -distribution assumes that observations are independent and follow a normal distribution: a sufficient, but not necessary, condition

Deviation from normality (heavier tails): test typically maintains type-I error control, but no longer has optimal power.

Options: use permutations; transform (e.g. ranks - Wilcoxon test)

Comments and Pitfalls

The proof that the t -statistic follows a t -distribution assumes that observations are independent and follow a normal distribution: a sufficient, but not necessary, condition

Deviation from normality (heavier tails): test typically maintains type-I error control, but no longer has optimal power.

Options: use permutations; transform (e.g. ranks - Wilcoxon test)

Deviation from independence: type-I error control is lost, p-values will likely be totally wrong (e.g., for positive correlation, too optimistic).

Comments and Pitfalls

The proof that the t -statistic follows a t -distribution assumes that observations are independent and follow a normal distribution: a sufficient, but not necessary, condition

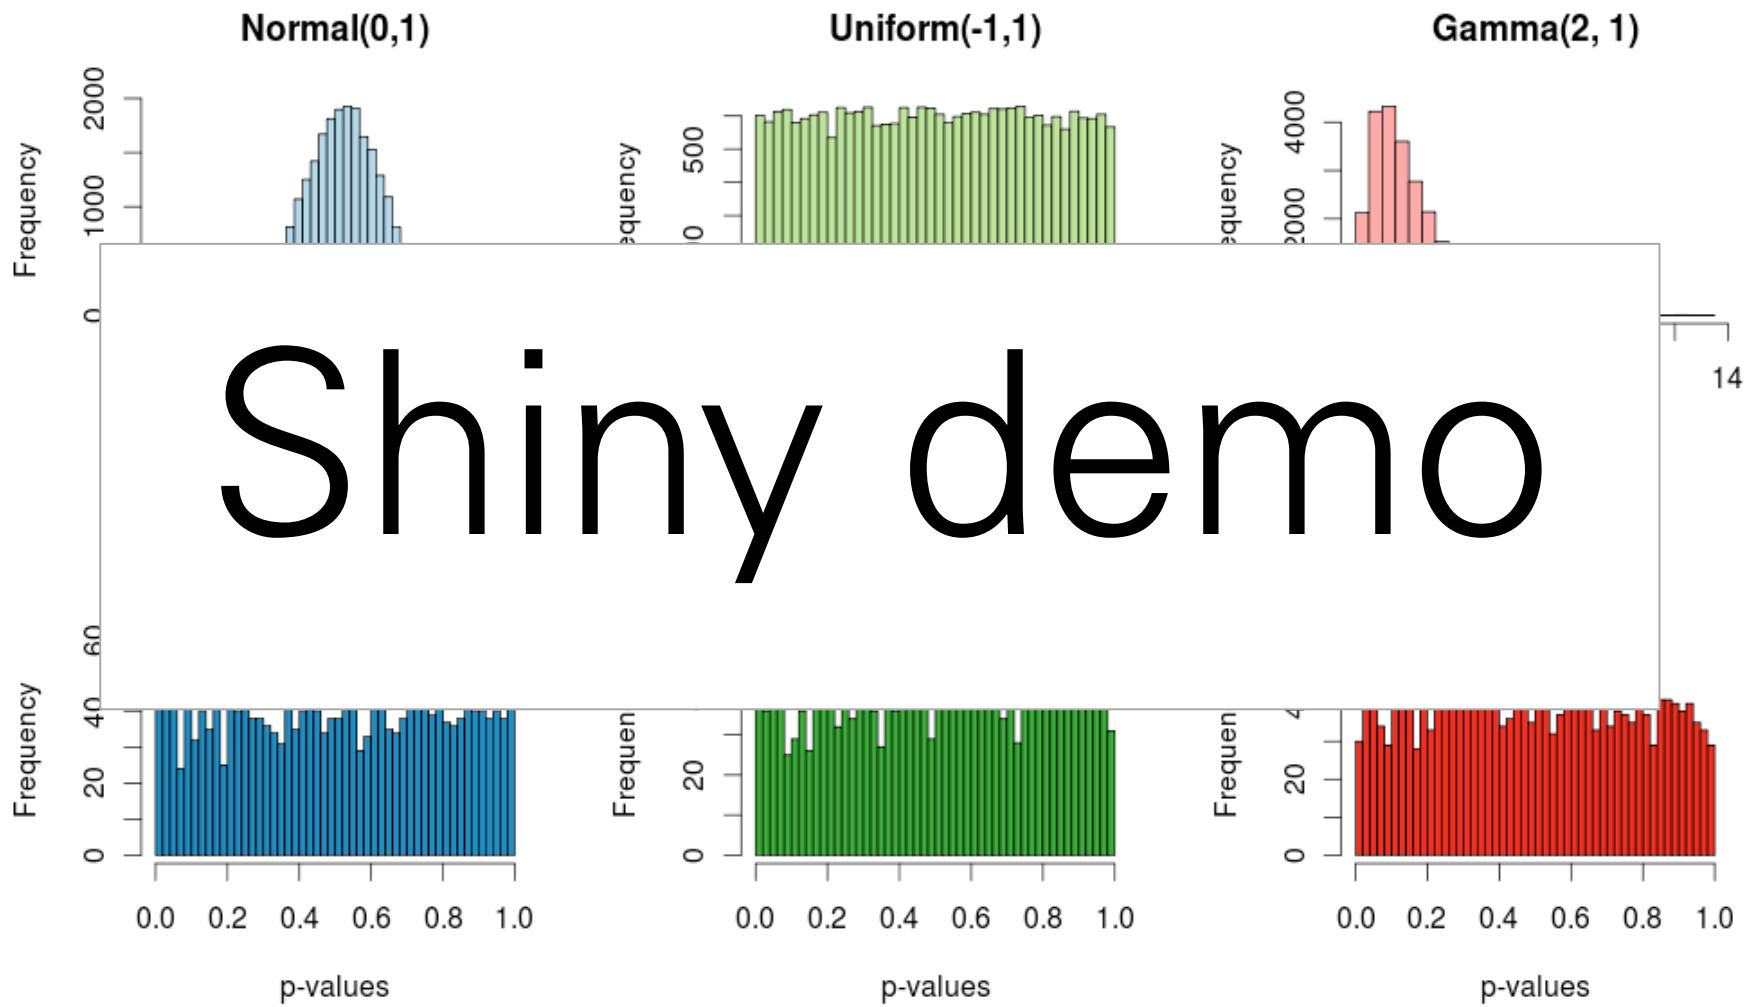
Deviation from normality (heavier tails): test typically maintains type-I error control, but no longer has optimal power.

Options: use permutations; transform (e.g. ranks - Wilcoxon test)

Deviation from independence: type-I error control is lost, p-values will likely be totally wrong (e.g., for positive correlation, too optimistic).

No easy options: ... try to model the dependence / remove it.... Efron's empirical null....

Effect of data distributions and dependence on t-test type I error control



Avoid Fallacy

The p-value is the probability that the data could happen, under the condition that the null hypothesis is true.

It is not the probability that the null hypothesis is true.

Absence of evidence ≠ evidence of absence



Recap: Single Hypothesis Testing

p-values are random variables: uniformly distributed if the null hypothesis is true - and should be close to zero if the alternative holds.

Note: We only observe one draw.

We prove something by disproving ('rejecting') the opposite (the null hypothesis). Reject = Discover.

Not rejecting does not prove the null hypothesis

Repeating the experiment (under the null): Around 5% of the times the p-value will be less than 0.05 by chance

All this reasoning is probabilistic. Testing & p-values are for rational decision making in uncertain contexts.

<u>P-VALUE</u>	<u>INTERPRETATION</u>
0.001	
0.01	
0.02	HIGHLY SIGNIFICANT
0.03	
0.04	
0.049	SIGNIFICANT
0.050	OH CRAP. REDO CALCULATIONS.
0.051	
0.06	ON THE EDGE OF SIGNIFICANCE
0.07	
0.08	HIGHLY SUGGESTIVE, SIGNIFICANT AT THE
0.09	P<0.10 LEVEL
0.099	HEY, LOOK AT
≥0.1	THIS INTERESTING SUBGROUP ANALYSIS

Limitations of p-value based hypothesis testing

Summarizing the data into one single number mushes together effect size and sample size ("with big enough sample size, anything becomes significant")

No place to take into account plausibility or 'prior' knowledge

Often, the 'null' is small (point-like), alternative is large (region-like)

The p-value has nothing or little to do with the probability of making a false discovery!

What is p-Value Hacking ?

On the same data, try different tests until one is significant

What is p-Value Hacking ?

On the same data, try different tests until one is significant

On the same data, try different hypotheses until one is significant (HARKing - hypothesizing after results are known)

What is p-Value Hacking ?

On the same data, try different tests until one is significant

On the same data, try different hypotheses until one is significant (HARKing - hypothesizing after results are known)

Moreover...:

retrospective data picking

'outlier' removal

the 5% threshold and publication bias

What is p-Value Hacking ?

On the same data, try different tests until one is significant

On the same data, try different hypotheses until one is significant (HARKing - hypothesizing after results are known)

Moreover....:

retrospective data picking

'outlier' removal

the 5% threshold and publication bias

The ASA's Statement on p-Values:
Context, Process, and Purpose
Ronald L. Wasserstein & Nicole A. Lazara
DOI: 10.1080/00031305.2016.1154108

What is p-Value Hacking ?

On the same data, try different tests until one is significant

On the same data, try different hypotheses until one is significant (HARKing - hypothesizing after results are known)

Moreover....:

retrospective data picking

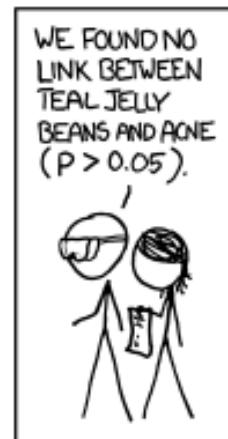
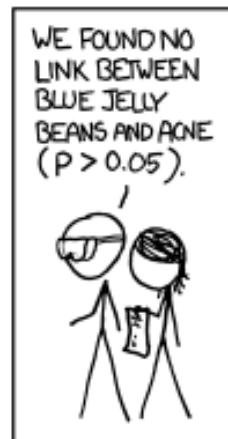
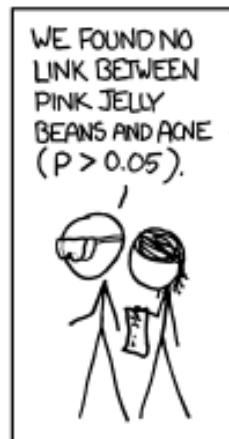
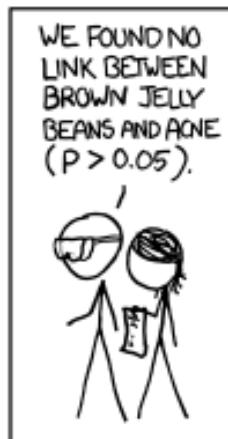
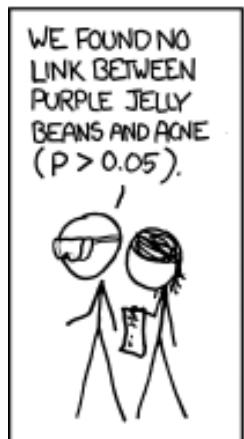
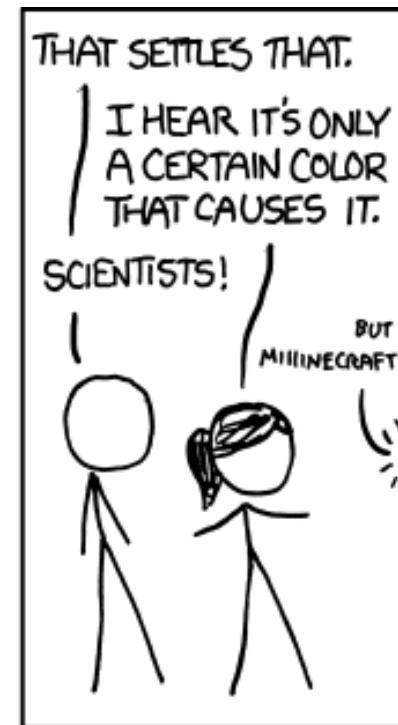
'outlier' removal

the 5% threshold and publication bias

The ASA's Statement on p-Values:
Context, Process, and Purpose
Ronald L. Wasserstein & Nicole A.
Lazara DOI:
[10.1080/00031305.2016.1154108](https://doi.org/10.1080/00031305.2016.1154108)

What can we do about this?

Multiple Testing



Multiple Testing

WE FOUND NO
LINK BETWEEN
SALMON JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
RED JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
TURQUOISE JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
MAGENTA JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
YELLOW JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
GREY JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
TAN JELLY
BEANS AND ACNE
($P > 0.05$).



WE FOUND NO
LINK BETWEEN
CYAN JELLY
BEANS AND ACNE
($P > 0.05$).



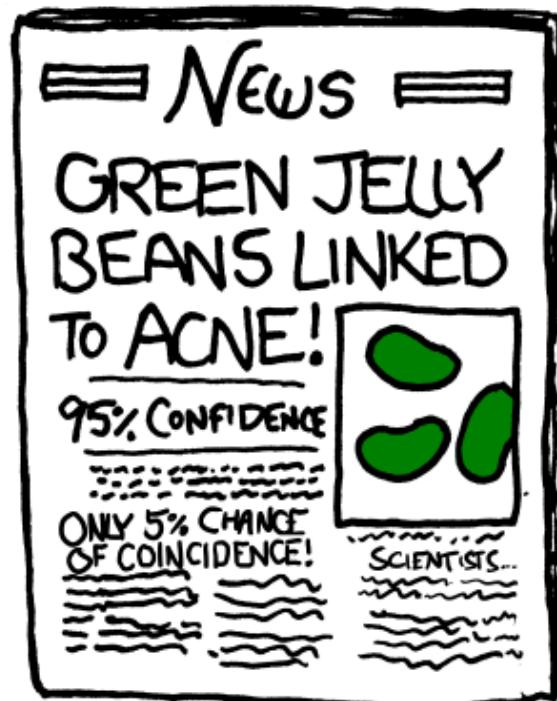
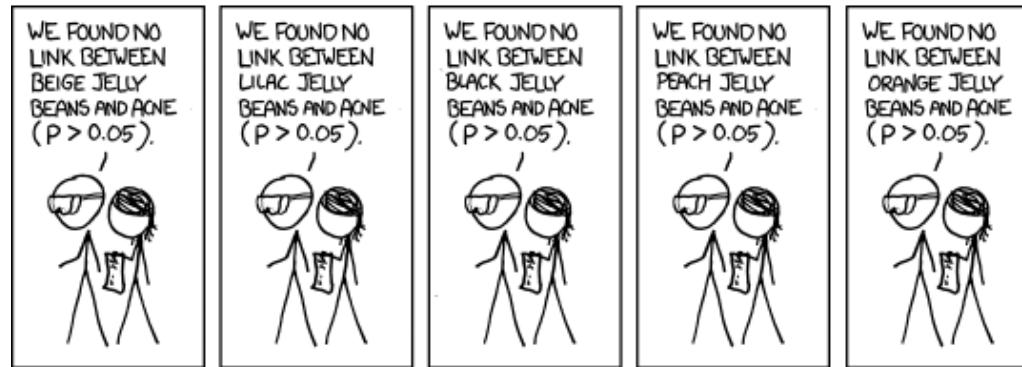
WE FOUND A
LINK BETWEEN
GREEN JELLY
BEANS AND ACNE
($P < 0.05$).



WE FOUND NO
LINK BETWEEN
MAUVE JELLY
BEANS AND ACNE
($P > 0.05$).



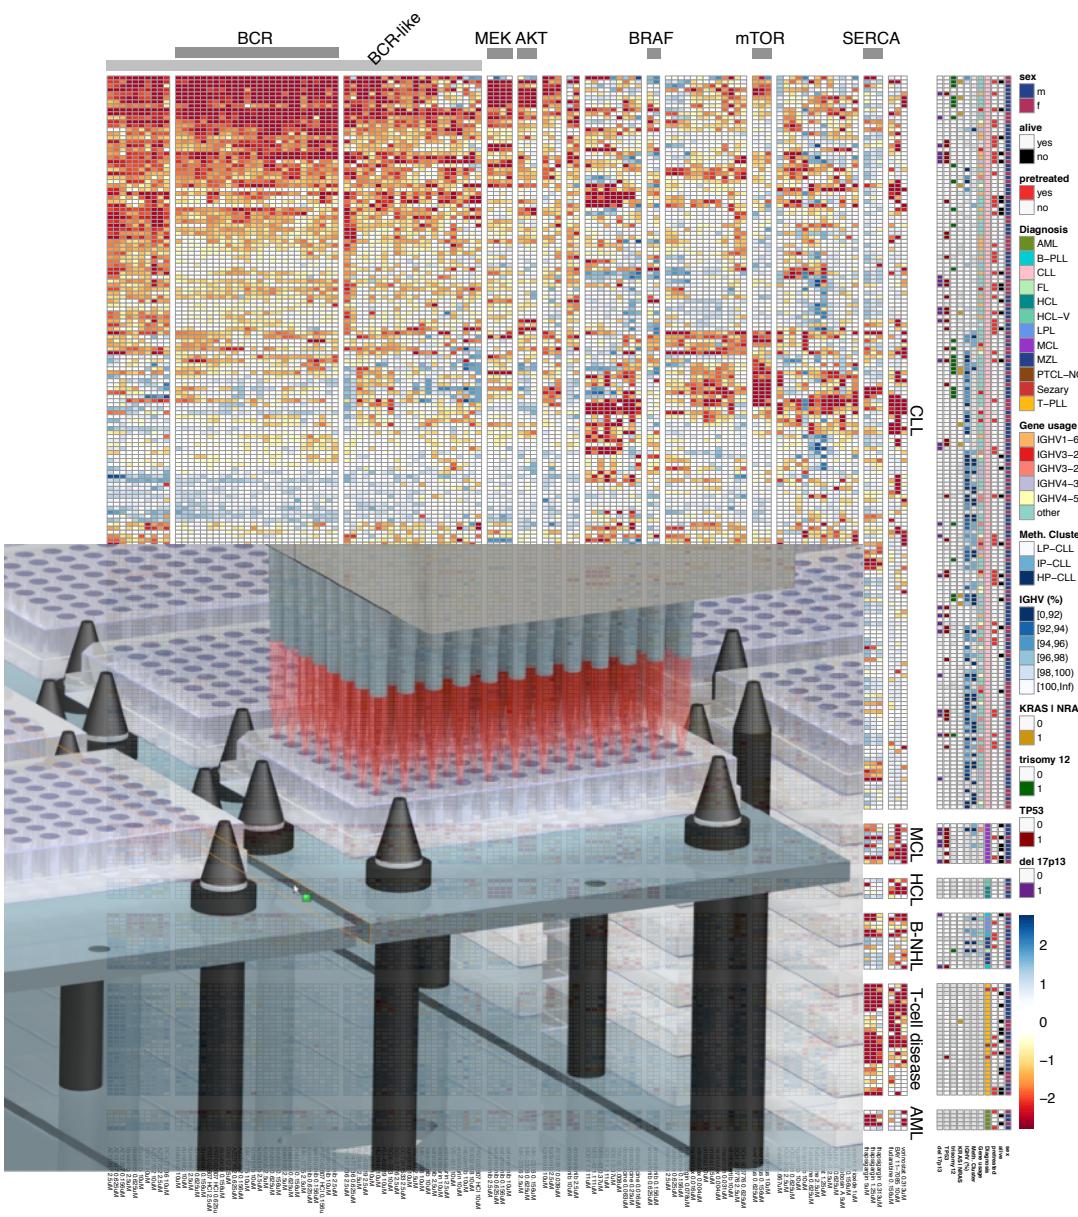
Multiple Testing



Multiple Testing

Many data analysis approaches in genomics employ item-by-item testing:

- Expression profiling
 - Differential microbiome analysis
 - Genetic or chemical compound screens
 - Genome-wide association studies
 - Proteomics
 - Variant calling
 - ...



False Positive Rate and False Discovery Rate

FPR: fraction of FP among all true negatives

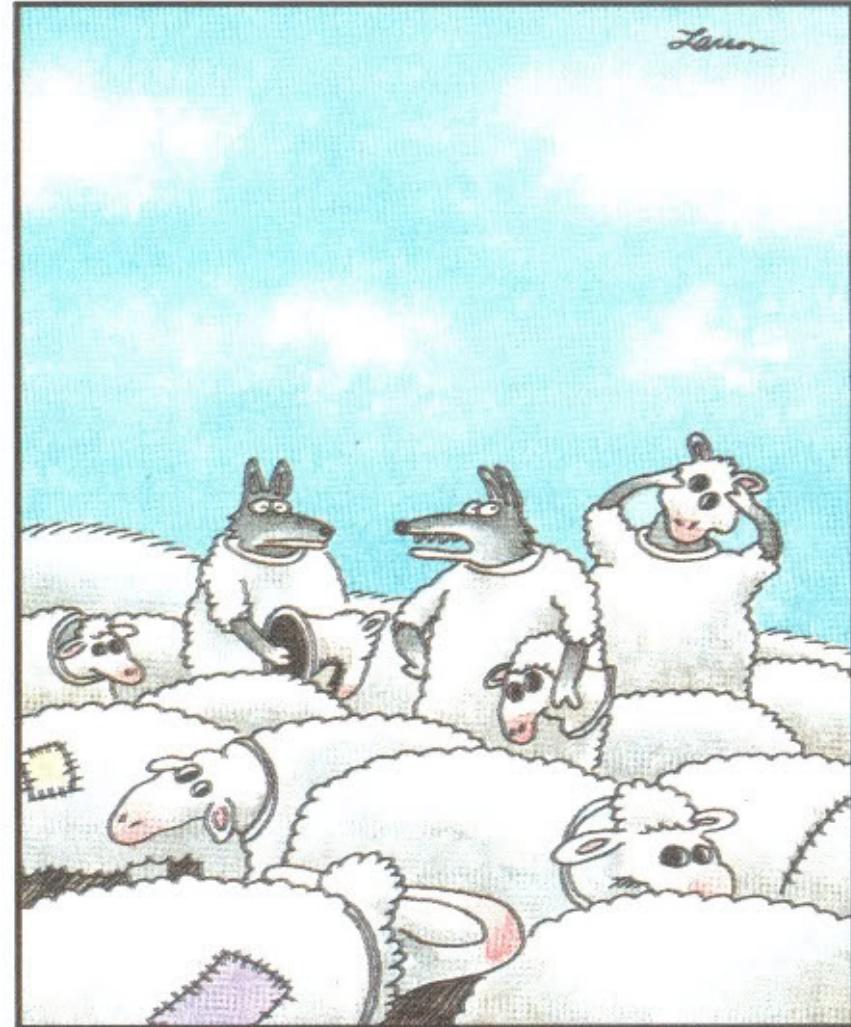
FDR: fraction of FP among hits called

Example:

20,000 genes, 500 are d.e.,
100 hits called, 10 of them wrong.

FPR: $10/19,500 \approx 0.05\%$

FDR: $10/100 = 10\%$



"Wait a minute! Isn't anyone here a real sheep?"

Experiment-Wide Type I Error Rates

Test vs Reality	Null Hypothesis is true	... is false	Total
Rejected	V	S	R
Not rejected	U	T	$m - R$
Total	m_0	$m - m_0$	m

- m : total number of hypotheses
- m_0 : number of null hypotheses
- V : number of false positives (a measure of type I error)

Family-wise error rate (FWER): The probability of one or more false positives, $P(V > 0)$. For large m_0 , this is difficult to keep small.

False discovery rate (FDR): The expected fraction of false positives among all discoveries, $E[V / \max \{R, 1\}]$.

NB: if $m_0 = m$, then FDR=FWER

The Multiple Testing Burden

When performing several tests, type I error goes up: for $\alpha = 0.05$ and n indep. tests, probability of no false positive result is

$$\underbrace{0.95 \cdot 0.95 \cdot \dots \cdot 0.95}_{n\text{-times}} \lll 0.95$$



Bonferroni Correction



For m tests, multiply each p -value with m .
Then see if anyone still remains below α .

The Multiple Testing Opportunity

DID THE SUN JUST EXPLODE?
(IT'S NIGHT, SO WE'RE NOT SURE.)

THIS NEUTRINO DETECTOR MEASURES
WHETHER THE SUN HAS GONE NOVA.

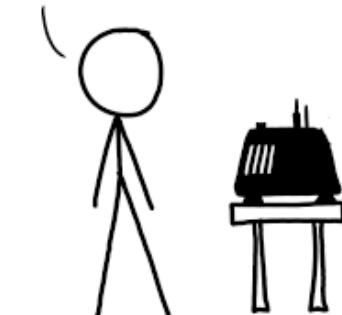
LET'S TRY.
DETECTOR! HAS THE
SUN GONE NOVA?

THEN, IT ROLLS TWO DICE. IF THEY
BOTH COME UP SIX, IT LIES TO US.
OTHERWISE, IT TELLS THE TRUTH.



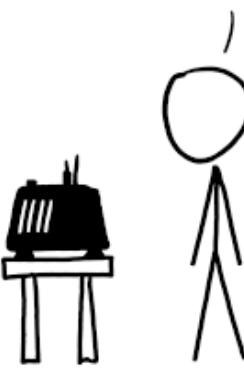
FREQUENTIST STATISTICIAN:

THE PROBABILITY OF THIS RESULT
HAPPENING BY CHANCE IS $\frac{1}{36} = 0.027$.
SINCE $p < 0.05$, I CONCLUDE
THAT THE SUN HAS EXPLODED.



BAYESIAN STATISTICIAN:

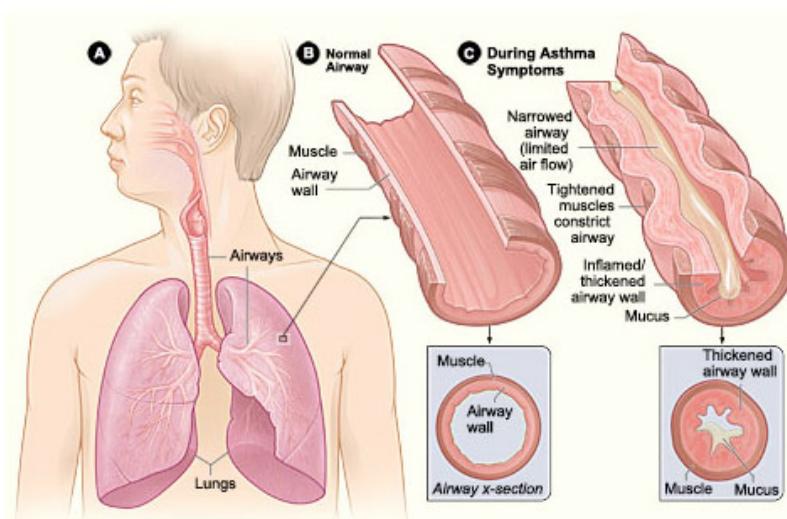
BET YOU \$50
IT HASN'T.



Data set 1: RNA-Seq

Transcriptome changes in four samples of primary human airway smooth muscle cells treated with dexamethasone, a synthetic glucocorticoid. 1 μM for 18 h.

cellline	dexamethasone
N61311	untrt
N61311	trt
N052611	untrt
N052611	trt
N080611	untrt
N080611	trt
N061011	untrt
N061011	trt



DESeq2 differential expression analysis:

gene i , sample j :

$$K_{ij} \sim \text{NB}(\text{mean} = \mu_{ij}, \text{dispersion} = \alpha_j)$$

$$\mu_{ij} = s_j q_{ij}$$

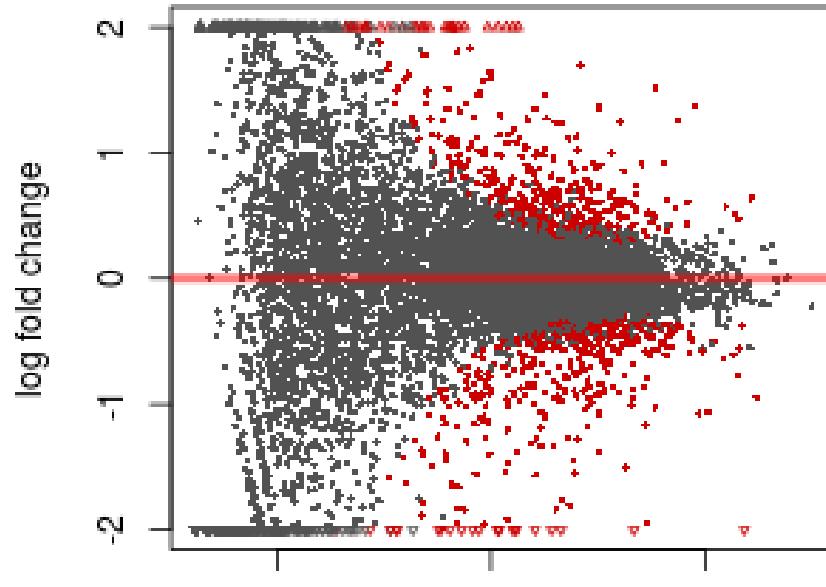
$$\log q_{ij} = \sum_r x_{jr} \beta_{rj}$$

design <- ~ cellline + dexamethasone

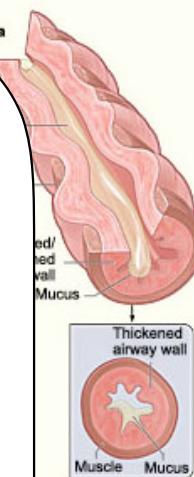
Data set 1: RNA-Seq

Transcript samples
smooth
dexamethasone
glucocorticoids

cellline
N61011
N61011
N61011
N05
N05
N08
N08
N061011
N061011 trt



design <- ~ cellline + dexamethasone



analysis:

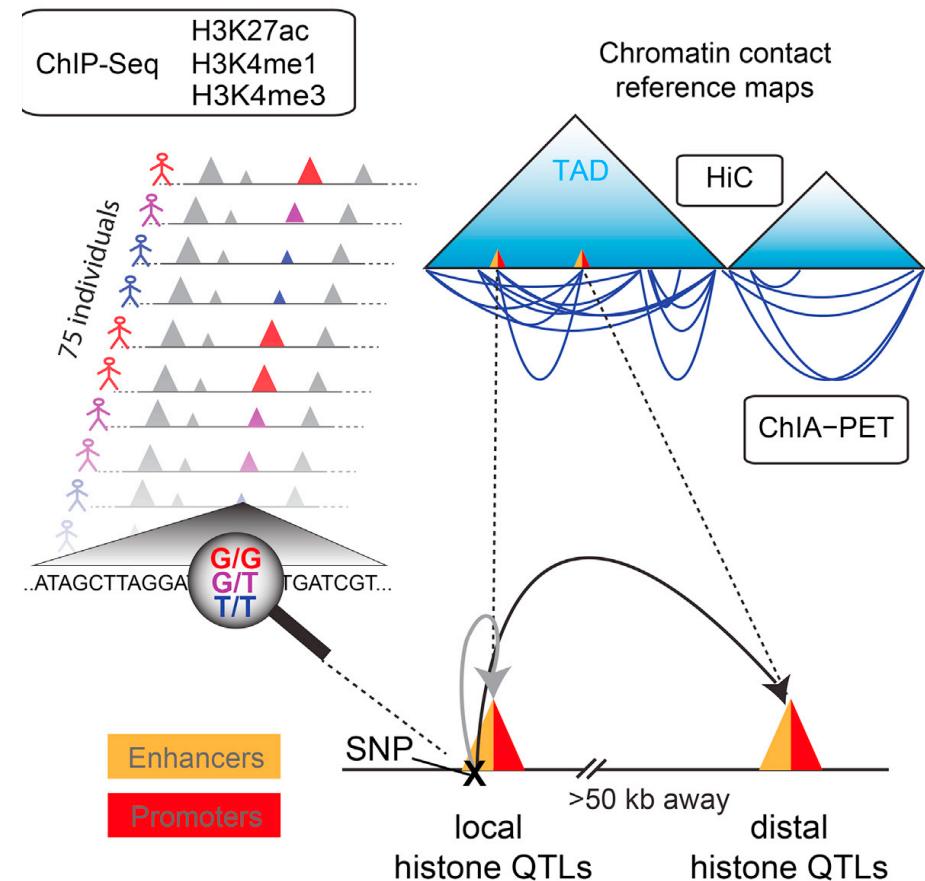
dispersion = α_j)

Data set 2: hQTL

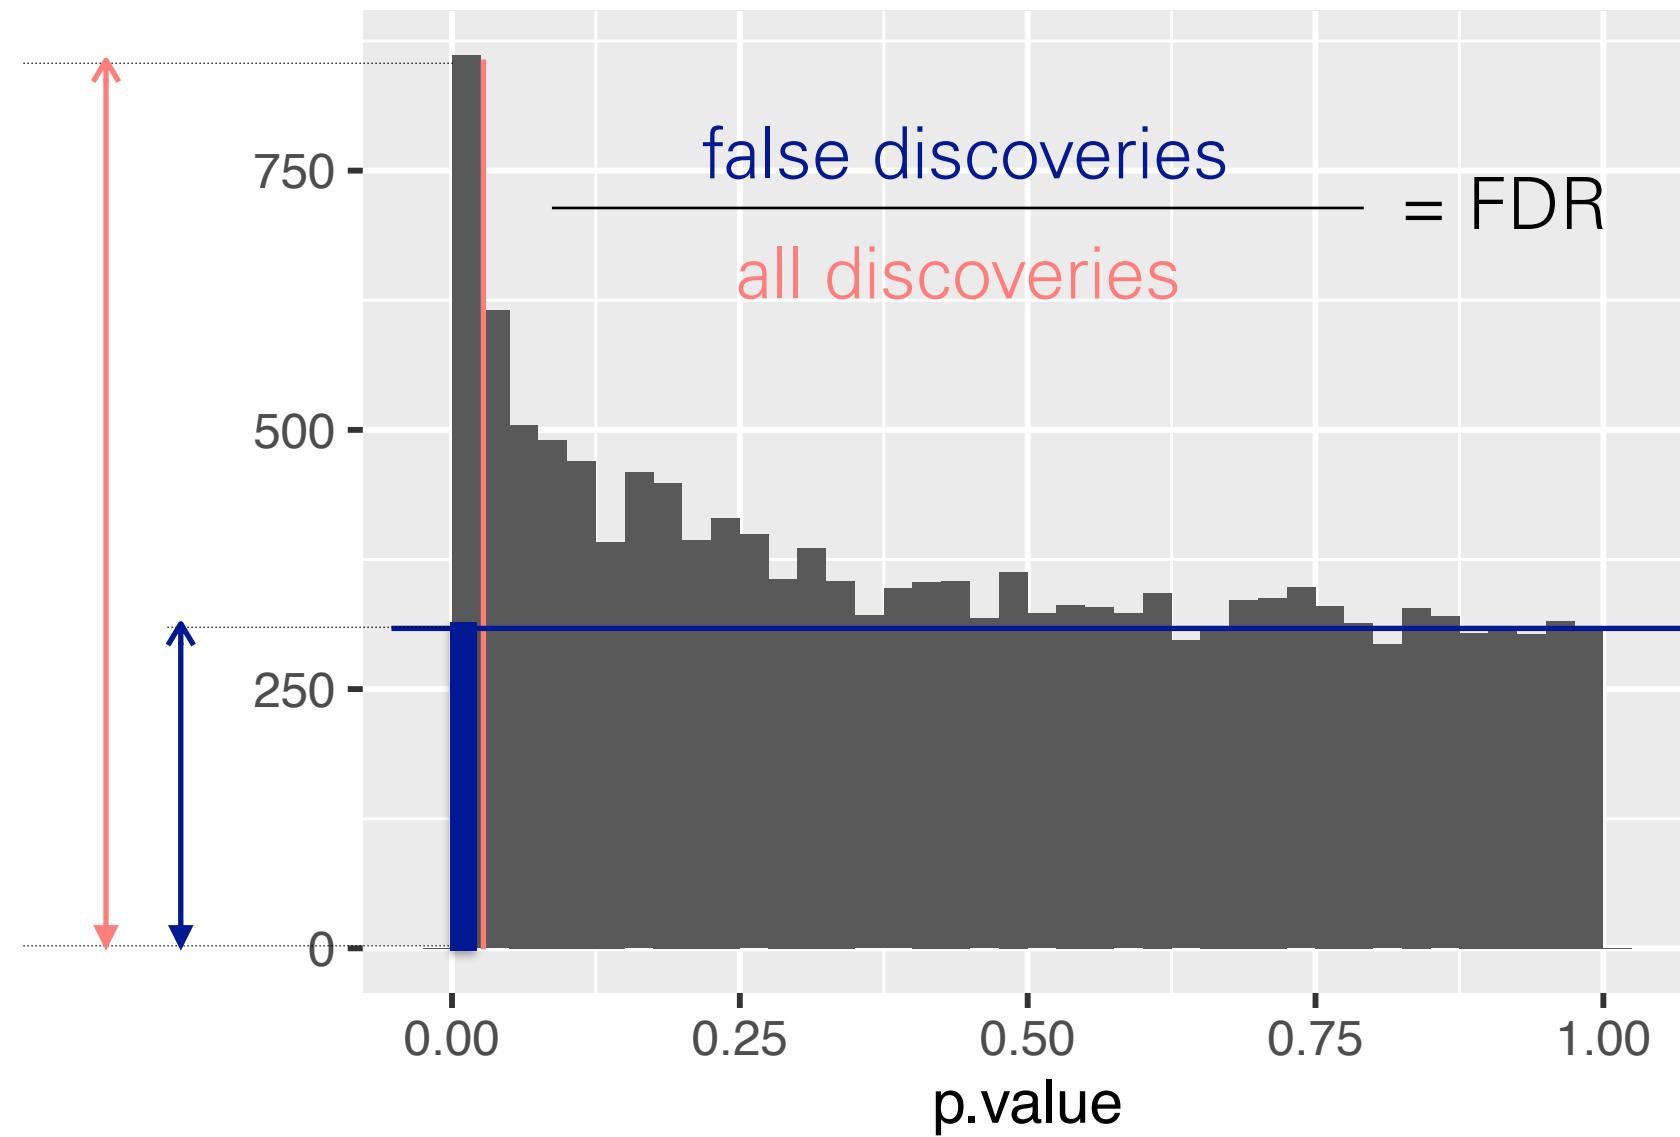
ChIP-seq for histone marks in lymphoblastoid cell lines from 75 sequenced individuals.

Local QTLs: find best-correlated SNP within 2kb of peak boundaries: 14,142 hQTLs, involving ~10% of all H3K27ac peaks (FDR=0.1, permutations)

Distal: distance cutoffs from 50 to 300 kb; also HiC

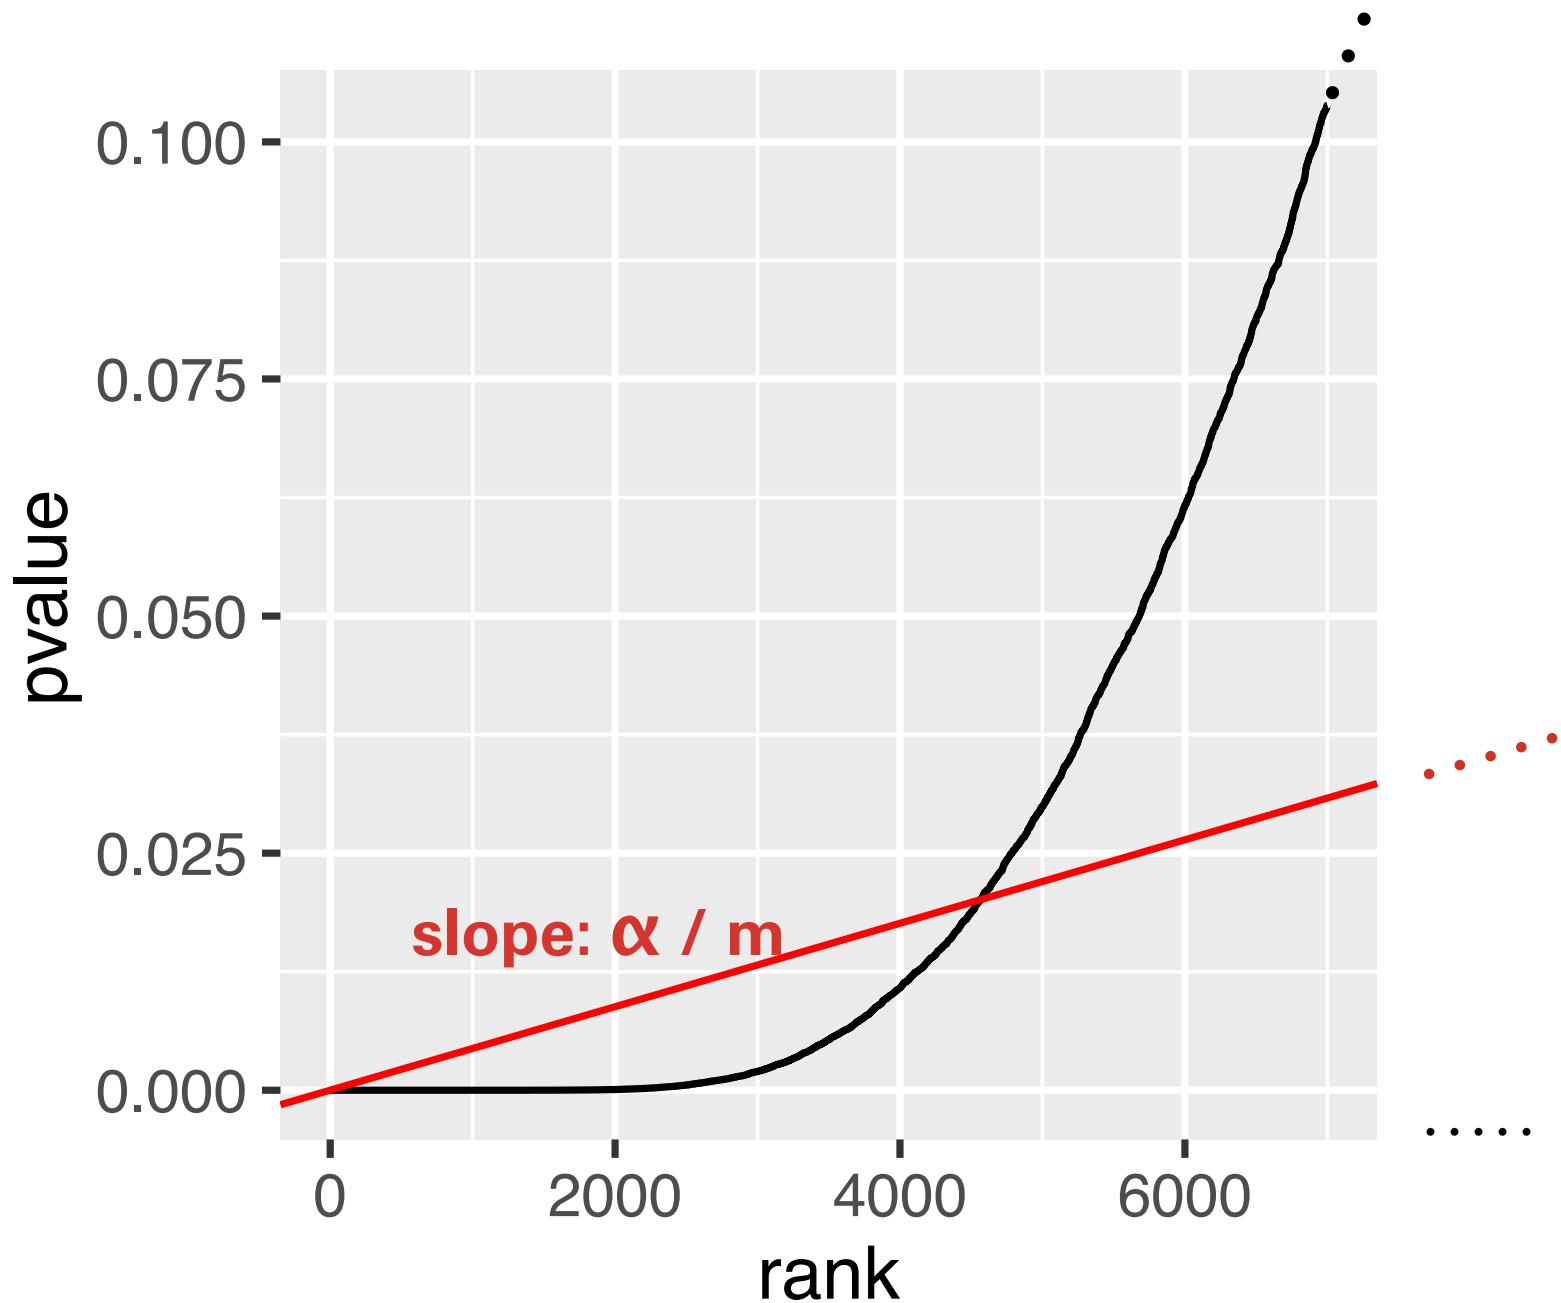


False Discovery Rate



Method of Benjamini & Hochberg (1995)

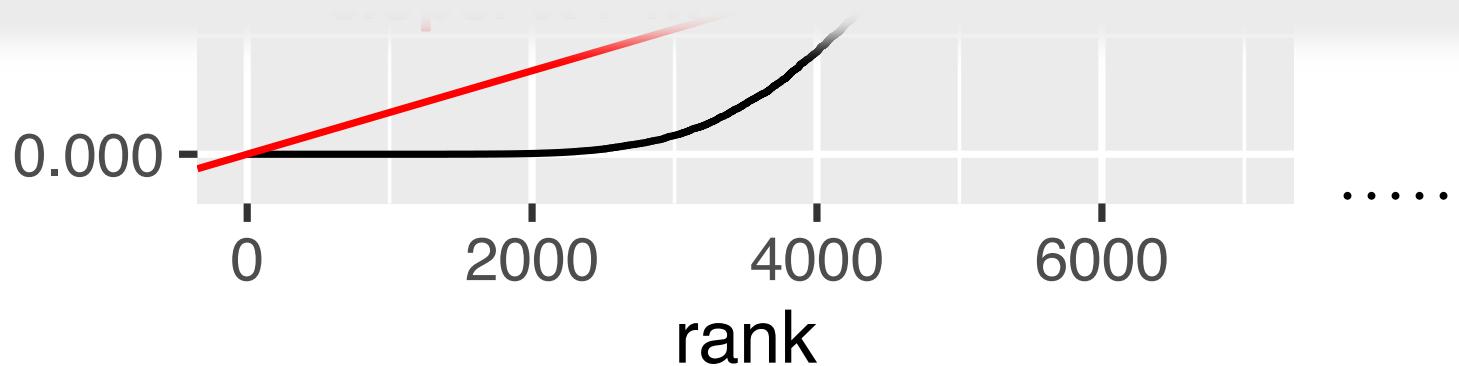
Method of Benjamini & Hochberg



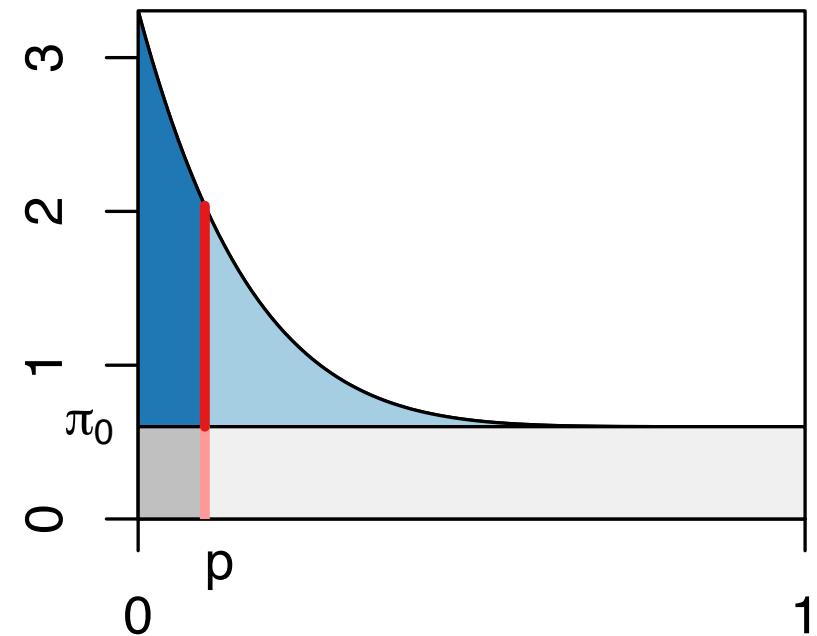
Method of Benjamini & Hochberg

```
BH = {  
    i <- length(p) :1  
    o <- order(p, decreasing = TRUE)  
    ro <- order(o)  
    pmin(1, cummin(n/i * p[o])) [ro]  
}
```

takes a list of p-values as input and returns a matched list of 'adjusted' p-values.



The Two-Groups Model and the Local False Discovery Rate



$$f(p) = \pi_0 + (1 - \pi_0)f_{\text{alt}}(p)$$

$$\text{fdr}(p) = \frac{\pi_0}{f(p)}$$

FDR: a set property. A single number that applies to a whole set of discoveries.

fdr: a local property. It applies to individual hypotheses.

Not all Hypothesis Tests are Created Equal

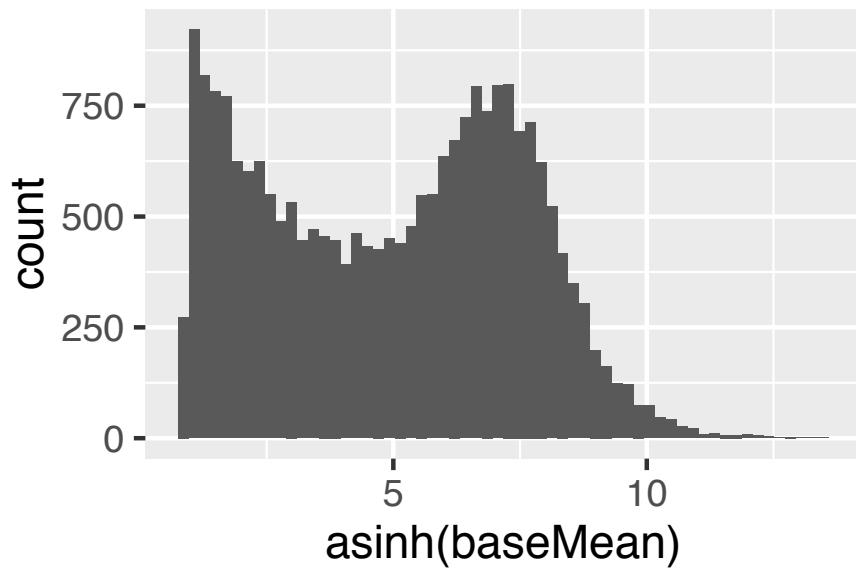
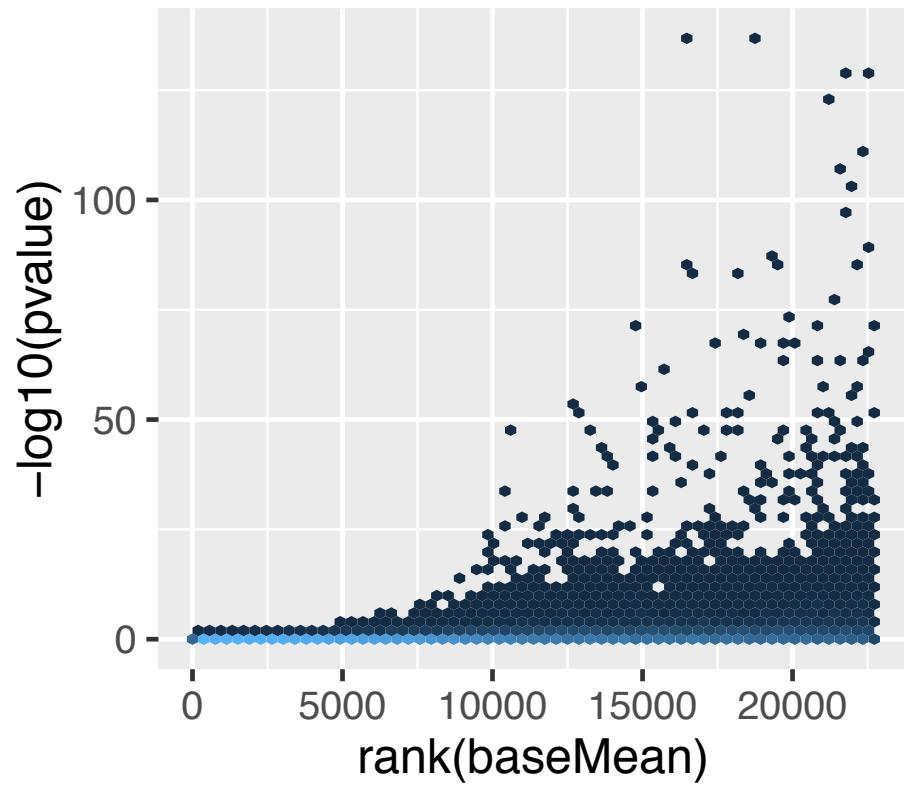


Figure 6.15: Histogram of baseMean . We see that it covers a large dynamic range, from close to 0 to around 3.3×10^5 .



Covariates - examples

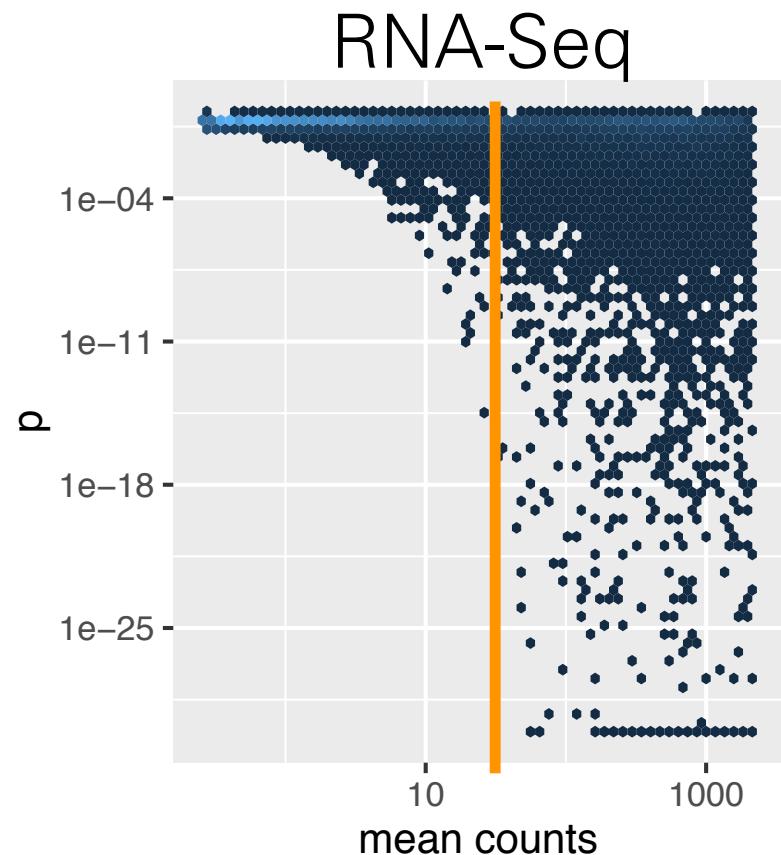
Application	Covariate
Differential RNA-Seq, ChIP-Seq, CLIP-seq, ...	(Normalized) mean of counts for each gene
eQTL analysis	SNP – gene distance
GWAS	Minor allele frequency
<i>t</i> -tests	Overall variance
Two-sided tests	Sign
All applications	Sample size; measures of signal-to-noise ratio

Independent Filtering

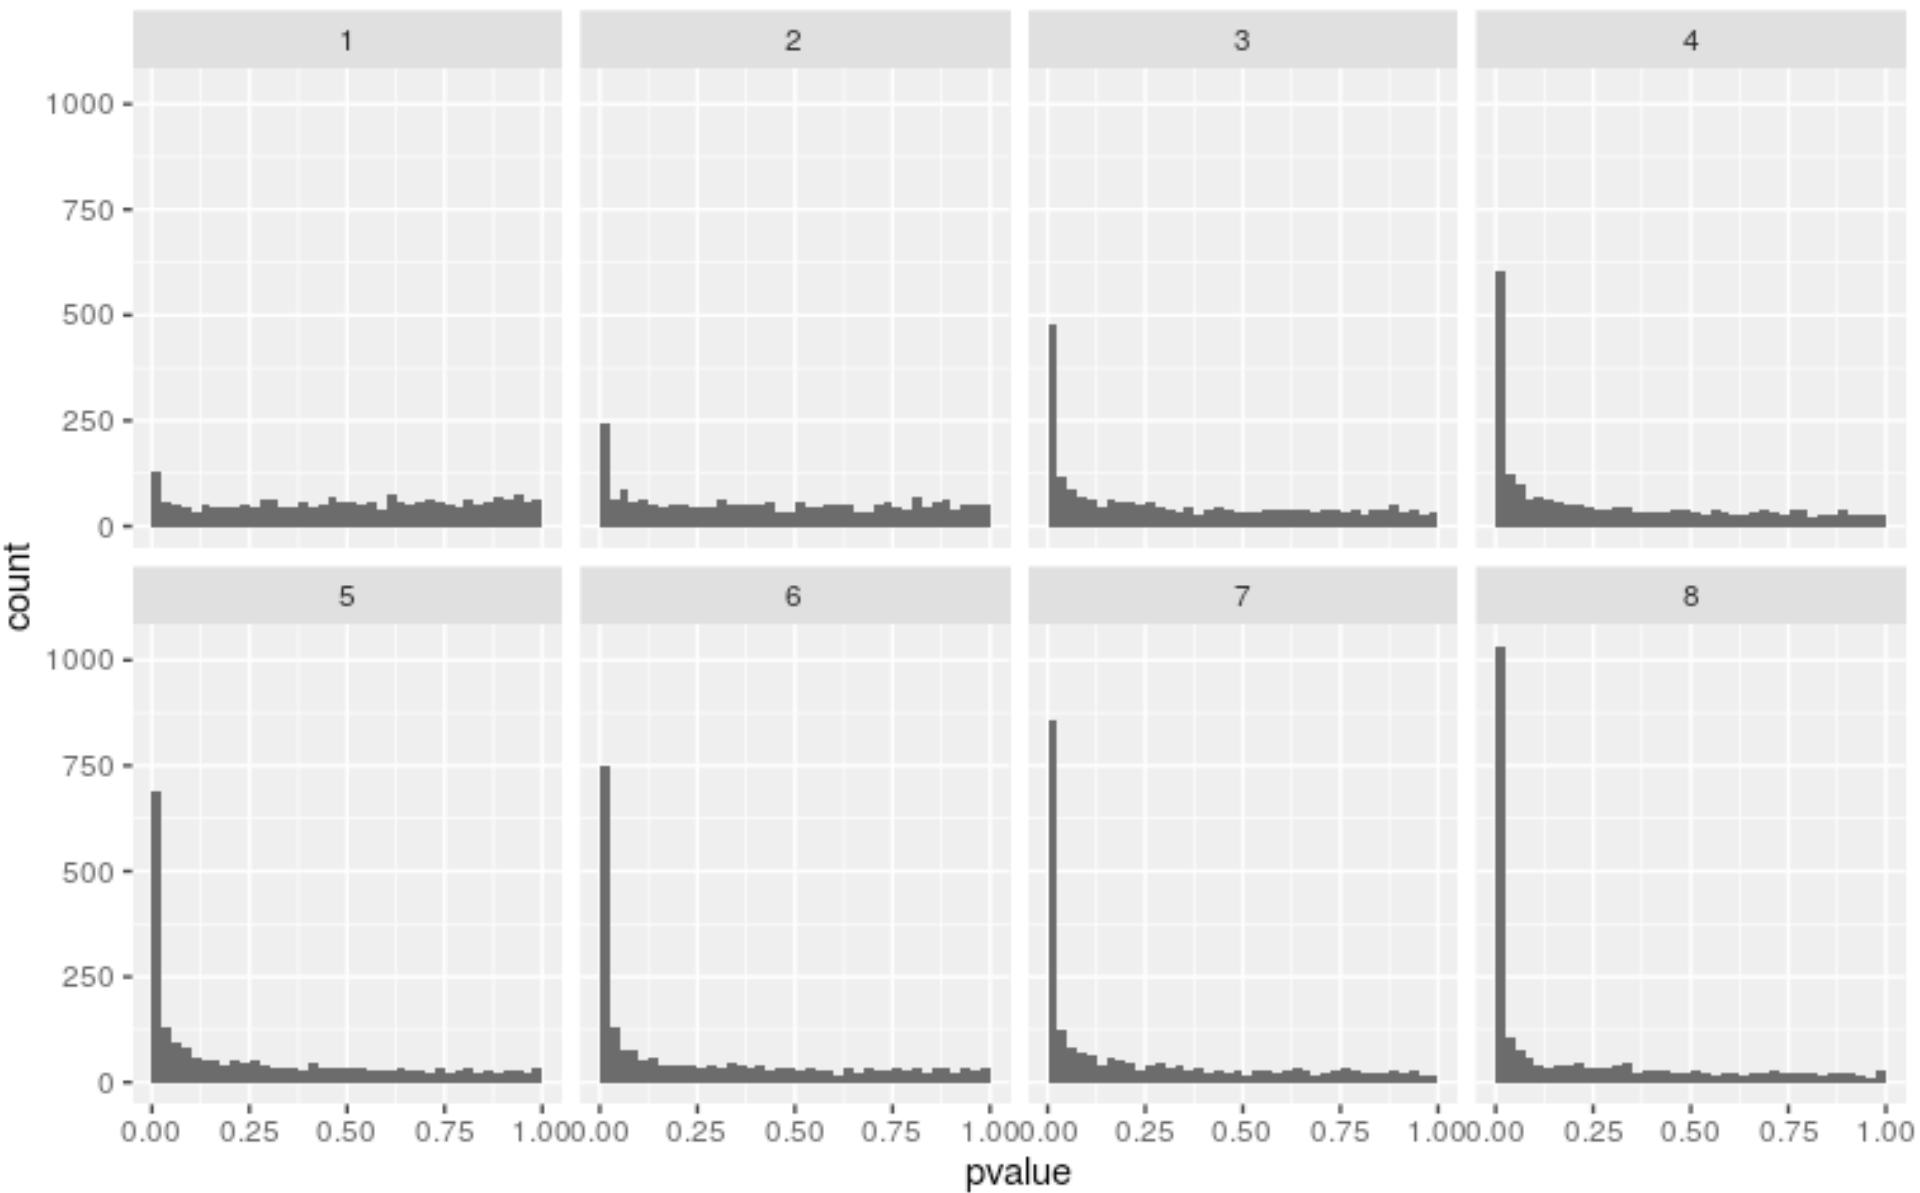
Two steps:

- All hypotheses H_i with $X_i < x$ get filtered.
- Apply BH to remaining hypotheses.

(Bourgon, Gentleman, Huber
PNAS 2010)



RNA-Seq p-value histogram stratified by average read count



Weighted Benjamini-Hochberg method

- Let $w_i \geq 0$ and $\frac{1}{m} \sum_{i=1}^m w_i = 1$ ("weight budget").
- Define $Q_i = P_i/w_i$.
- Apply BH to Q_i instead of P_i .
- Proven Type-I error (FDR) control (Genovese, Roeder, Wasserman *Biometrika* 2006).
- If $w_i > 1$, then H_i is easier to reject.
- $Q_i \leq t \Leftrightarrow P_i \leq w_i t =: t_i$

Weighted Benjamini-Hochberg method

- Let $w_i \geq 0$ and $\frac{1}{m} \sum_{i=1}^m w_i = 1$ ("weight budget").
- Define $Q_i = P_i/w_i$.
- Apply BH to Q_i instead of P_i .
- Proven Theory by Benjamini, Krieger, and Yekutieli; Wasserman and Altonen
- If $w_i > 1$, then $Q_i = P_i$.
- $Q_i \leq t \Leftrightarrow P_i \leq t w_i$



Weighted Benjamini-Hochberg method

- Let $w_i \geq 0$ and $\frac{1}{m} \sum_{i=1}^m w_i = 1$.
Problem: how to know the “budget”).
- Define $Q_i = P(w_i)$.
- Apply BH procedure.
- Proven by Benjamini, Krieger, Wasserman (1997).
- If $w_i > 1$, then $Q_i < t$.
- $Q_i \leq t \Leftrightarrow$

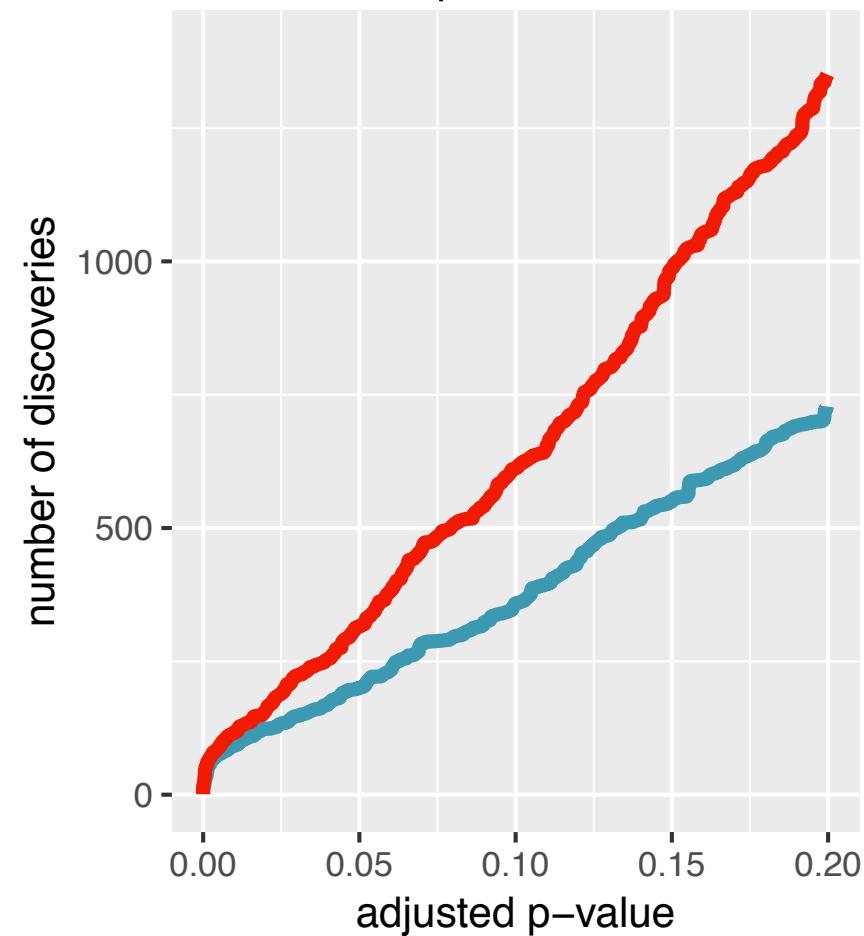


Independent hypothesis weighting (IHW): basic idea

- Stratify the tests into G bins, by covariate X
- Choose α
- For each possible weight vector $\mathbf{w} = (w_1, \dots, w_G)$ apply weighted BH procedure. Choose \mathbf{w} that maximizes the number of rejections at level α .
- Report the result with the optimal weight vector \mathbf{w}^* .

RNA-Seq example (DESeq2)

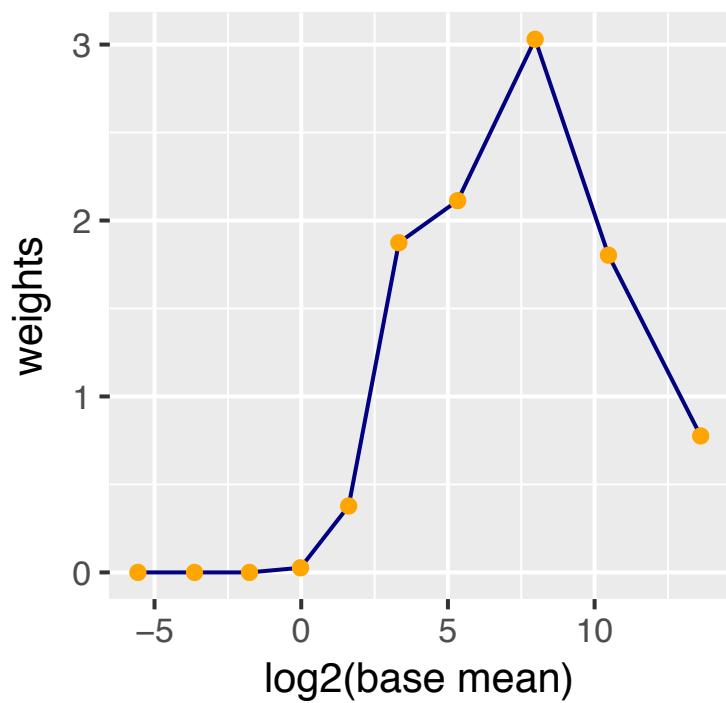
power



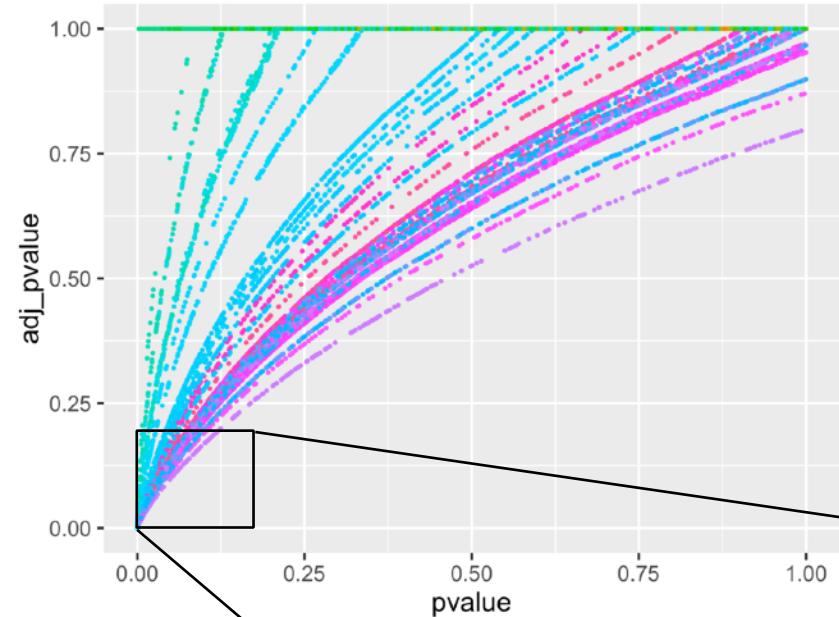
method

- BH
- IHW

weights

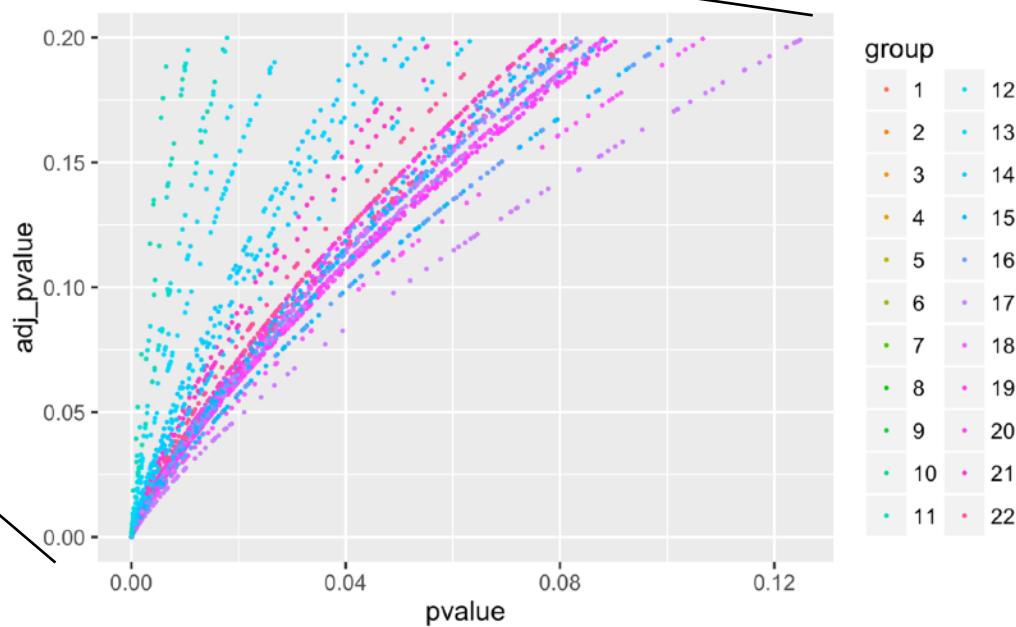


Ranking is not monotonous in raw p-values



group

1	12
2	13
3	14
4	15
5	16
6	17
7	18
8	19
9	20
10	21
11	22



group

1	12
2	13
3	14
4	15
5	16
6	17
7	18
8	19
9	20
10	21
11	22

Which estimators to use?



Nikos Ignatiadis

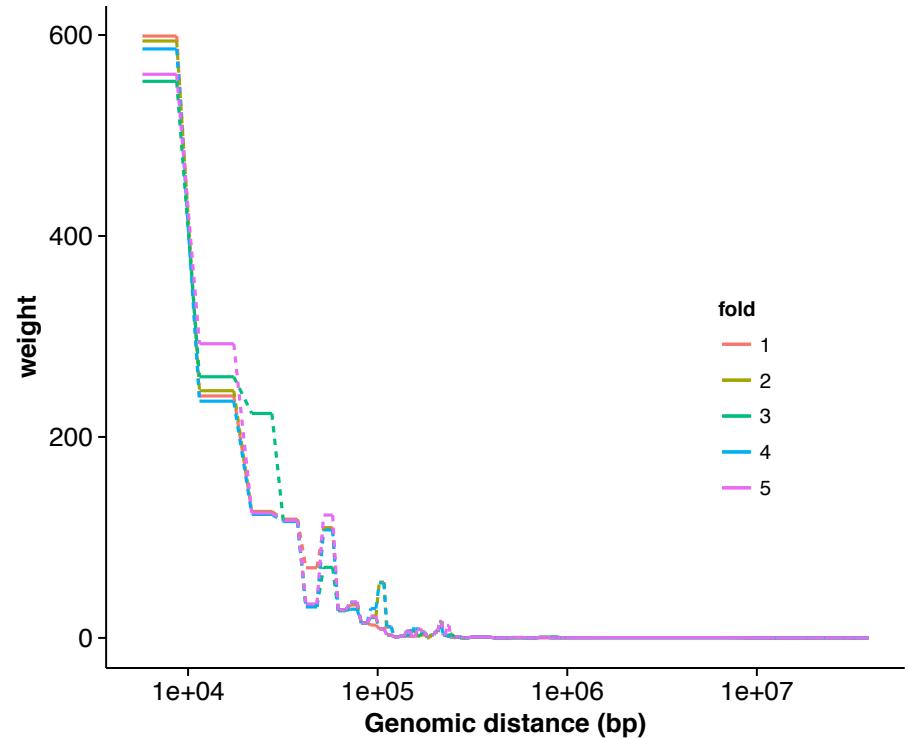
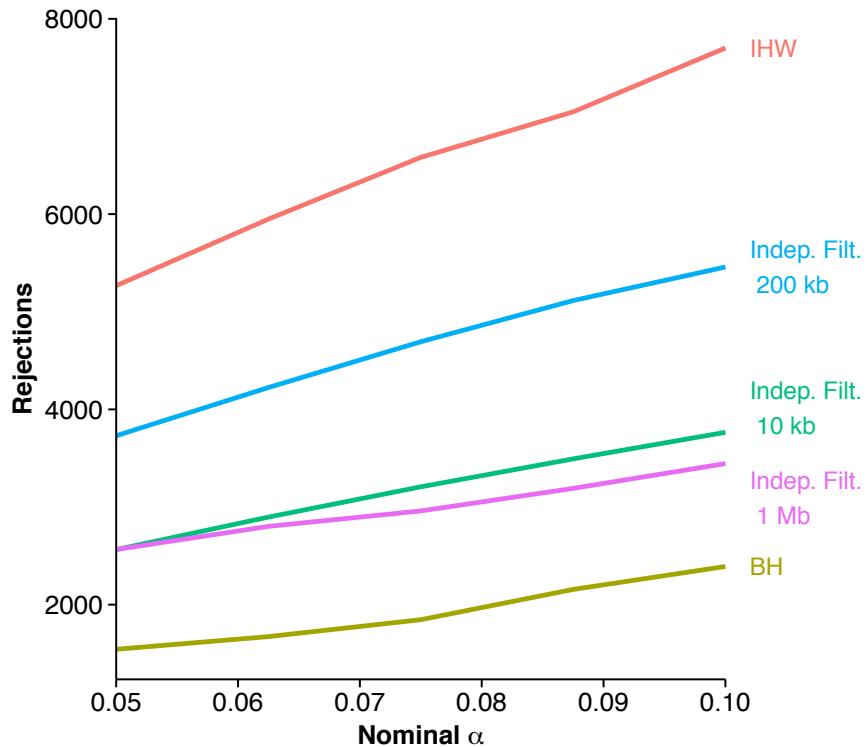
Hypothesis splitting: to ensure independence of 'training set' and 'test set'.
By explicit knowledge; or randomized

Regularisation:

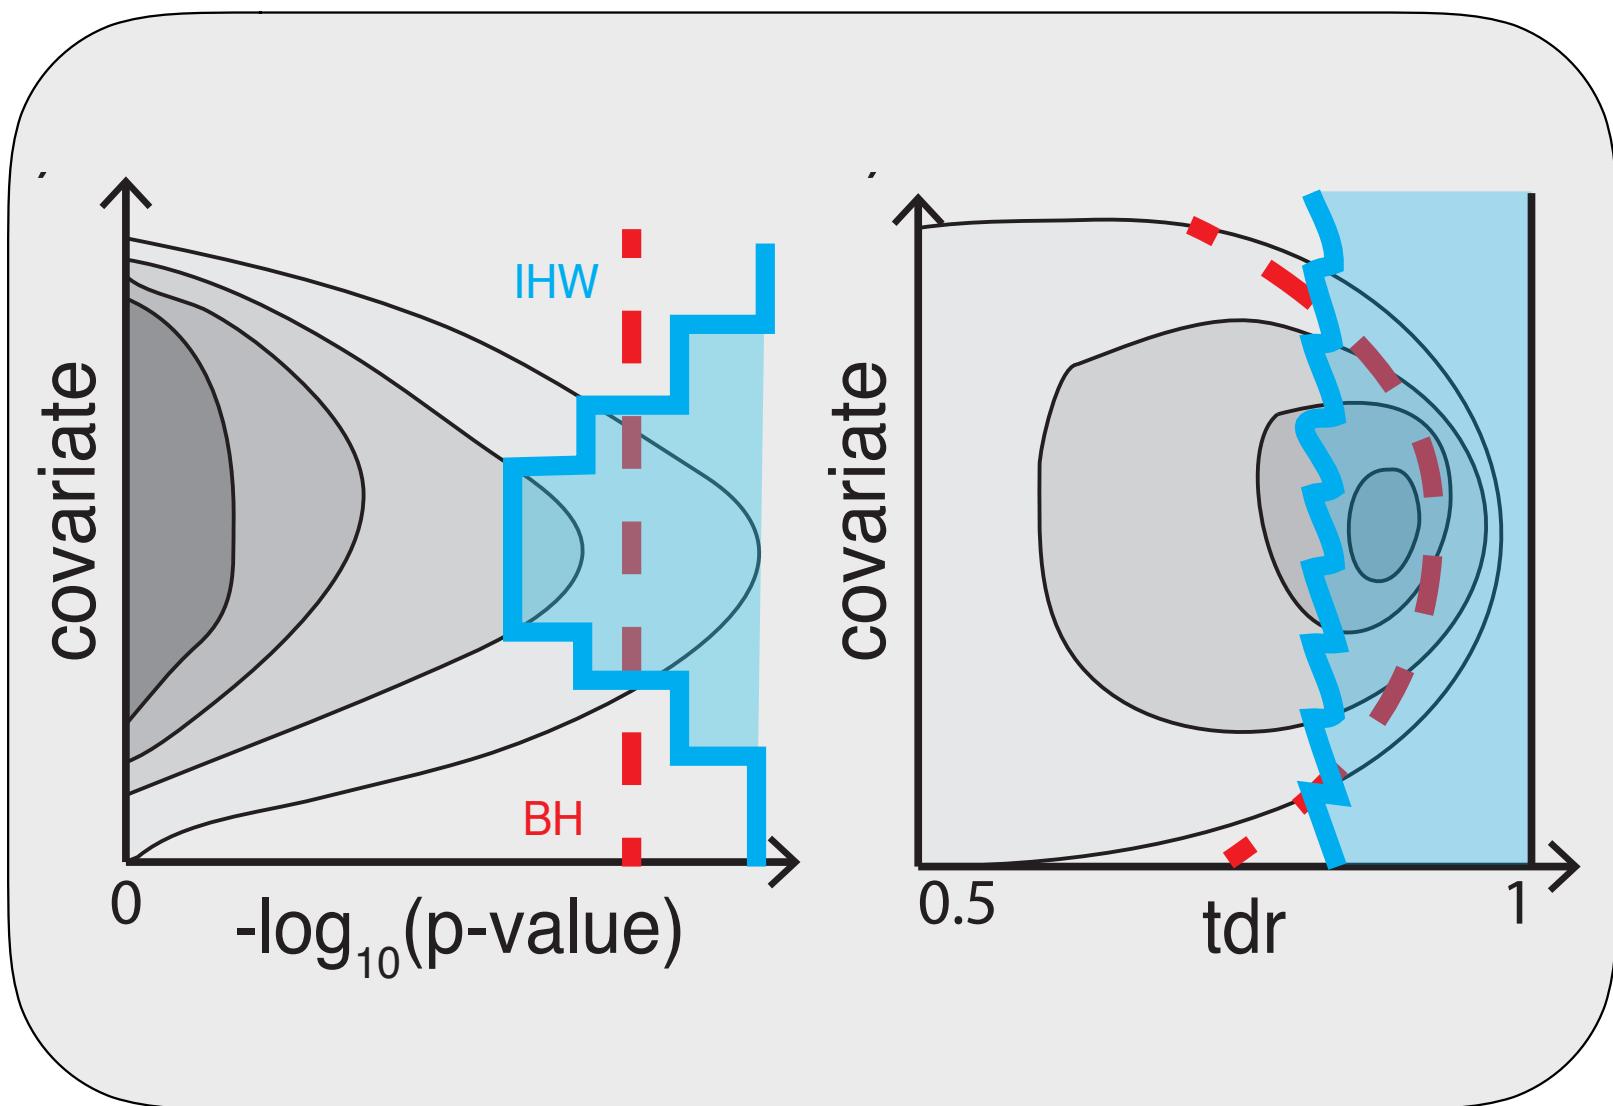
- for ordered covariate: $\sum_g |w_g - w_{g-1}| \leq \lambda$
- for categorical covariate: $\sum_g |w_g - 1| \leq \lambda$

Convex relaxation: for weight optimisation (only), replace ECDFs of the p-values with Grenander estimators (least concave majorant of the ECDF)

Histone-QTL example (H3K27ac)



2D decision boundaries



Summary

- Multiple testing is not a problem but an opportunity
- Heterogeneity across tests
- Informative covariates are often apparent to domain scientists
 - independent of test statistic under the null
 - informative on π_1 , F_{alt}
- Data-driven weighting
- Scales well to millions of hypotheses
- Controlling ‘overoptimism’



Nikos Ignatiadis

A promotional poster for the James Bond film "The World Is Not Enough". Pierce Brosnan plays James Bond, dressed in a classic grey suit, white shirt, and patterned tie. He stands in a dramatic pose against a background of intense orange and yellow flames, with smoke rising behind him. The title "The World Is Not Enough" is written in a large, white, serif font, with "The" and "Is Not" stacked above "Enough". Below the title, the iconic "007" logo is displayed in its signature gold and black color scheme.

The p-value ***is Not Enough***

Availability

Bioconductor - IHW (devel...)

https://www.bioconductor.org/packages/devel/t/

Search

Home Install Help Developers About

Search:

Home » Bioconductor 3.3 » Software Packages » IHW (development version)

IHW

platforms all downloads available posts 0 in Bioc devel only
build ok commits 0.17 test coverage unknown

This is the **development** version of IHW; to use it, please install the [devel version](#) of Bioconductor.

Independent Hypothesis Weighting

Bioconductor version: Development (3.3)

Independent hypothesis weighting (IHW) is a multiple testing procedure that increases power compared to the method of Benjamini and Hochberg by assigning data-driven weights to each hypothesis. The input to IHW is a two-column table of p-values and covariates. The covariate can be any continuous-valued or categorical variable that is thought to be informative on the statistical properties of each hypothesis test, while it is independent of the p-value under the null hypothesis.

Author: Nikos Ignatiadis [aut, cre]

Maintainer: Nikos Ignatiadis <nikos.ignatiadis01 at gmail.com>

Citation (from within R, enter `citation("IHW")`):

Ignatiadis N, Klaus B, Zaugg J and Huber W (2015). "Data-driven hypothesis weighting increases detection power in big data analytics." *bioRxiv*.

Installation

To install this package, start R and enter:

Documentation »

Bioconductor

- Package [vignettes](#) and manuals.
- [Workflows](#) for learning and use.
- [Course and conference](#) material.
- [Videos](#).
- Community [resources](#) and [tutorials](#).

R / [CRAN](#) packages and [documentation](#)

Support »

Please read the [posting guide](#). Post questions about Bioconductor to one of the following locations:

- [Support site](#) - for questions about Bioconductor packages
- [BioC-devel](#) mailing list - for package developers

Papers:

Nature Methods, June 2016

arXiv, January 2017

Further reading

www.huber.embl.de/msmb
Chapter 6

The screenshot shows a web browser window with multiple tabs open, including one for the EMBL calendar and another for the CSAMA 2018 programme. The main content is the Cambridge University Press Academic website.

Header: The header features the Cambridge University Press logo, navigation links for Academic, Cambridge English, Education, Bibles, Digital Products, About Us, and Careers, and a language selection for Italy. A search bar, a cart icon (Cart(0)), and user links for Sign in and Register are also present.

Breadcrumbs: Home / Academic / Statistics and probability / Statistics for life sciences, medicine and health /

Product Page: The page displays the book "Modern Statistics for Modern Biology" by Susan Holmes and Wolfgang Huber. It includes the Cambridge University Press logo, a "TEXTBOOK" badge, author information, publication details (Planned for November 2018), availability (Not yet published - available from November 2018), format (Hardback), ISBN (9781108427029), and a "Rate & review" button. The price is listed as c.£ 45.00 for Hardback, with "Pre-order" and "Add to wishlist" buttons. A "Request inspection copy" section allows lecturers to request a copy for inspection.

Footer: The footer includes links for Description, Contents, Resources, Courses, and About the Authors.